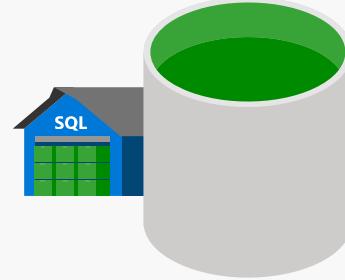


Azure SQL Data Warehouse



Contents

- Overview
- Architecture
- Scaling
- Schema Design
- Loading
- Querying
- High-Availability
- Backup & Disaster Recovery
- Monitoring, Management
- Security
- Sizing
- Migration
- Best Practices
- Pricing

Azure SQL Data Warehouse

Overview

Azure SQL Data Warehouse

Enables
transformative insights

Near real-time analytics
over structured and
unstructured data

Use tools you know and
love; Excel Power BI,
Tableau, Qlik or 20 others

Seamless compatibility with
Azure ingestion, machine
learning, and Big Data
services



Scales in minutes
with more freedom

Easily scale your
workload in minutes

Decoupled storage and
compute; pay for only
the compute you use

Innovative “Pause”
feature to stop compute
and save even more



Helps secure
and protect

Multiple layers of data
protection

Encryption at rest, in
motion or in use

Tools to monitor data
security

Protect user access to
data with Azure Active
Directory



Gets you up
& running quickly

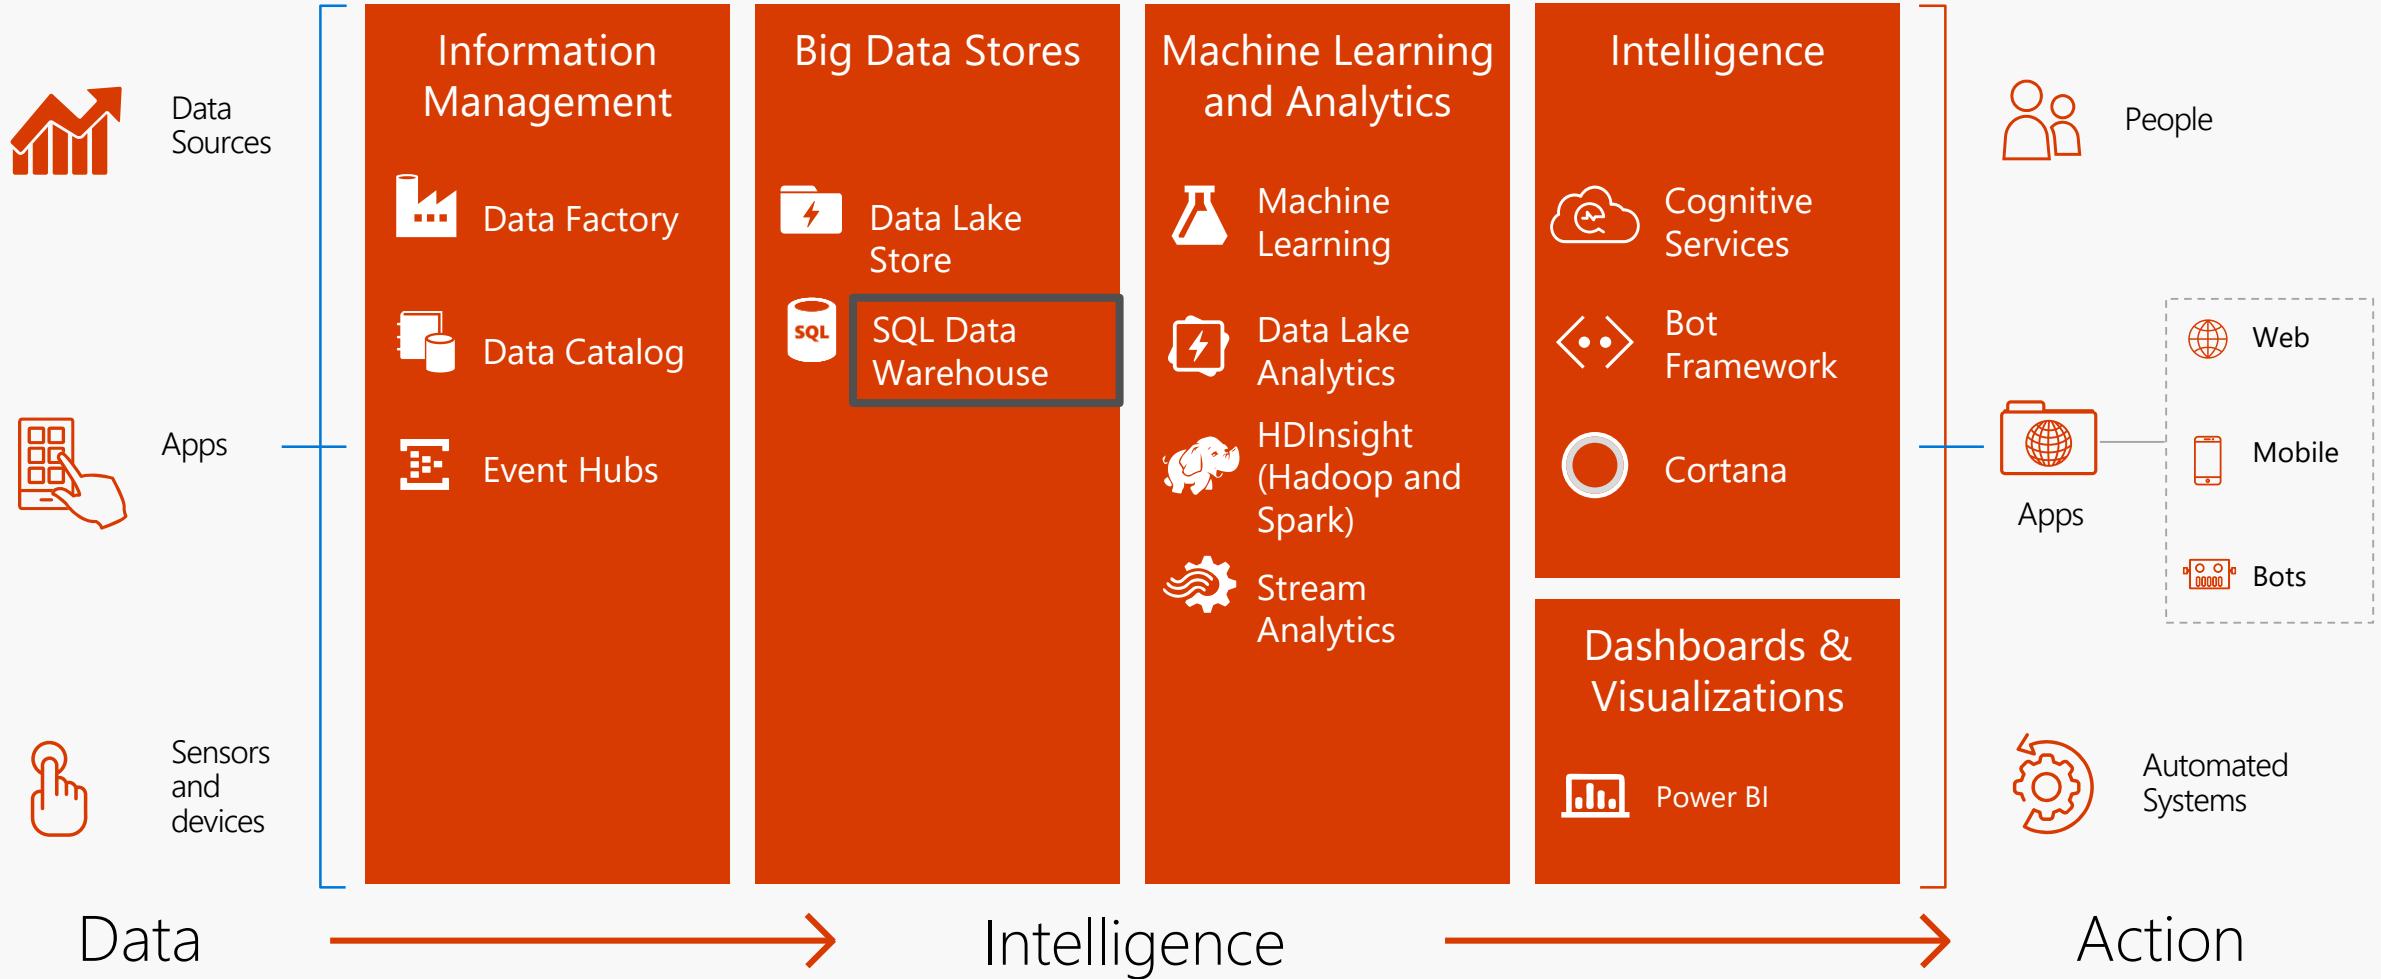
A single SQL-based view
across data sources

Seamlessly ingest data
from on-premises, cloud,
devices, and apps

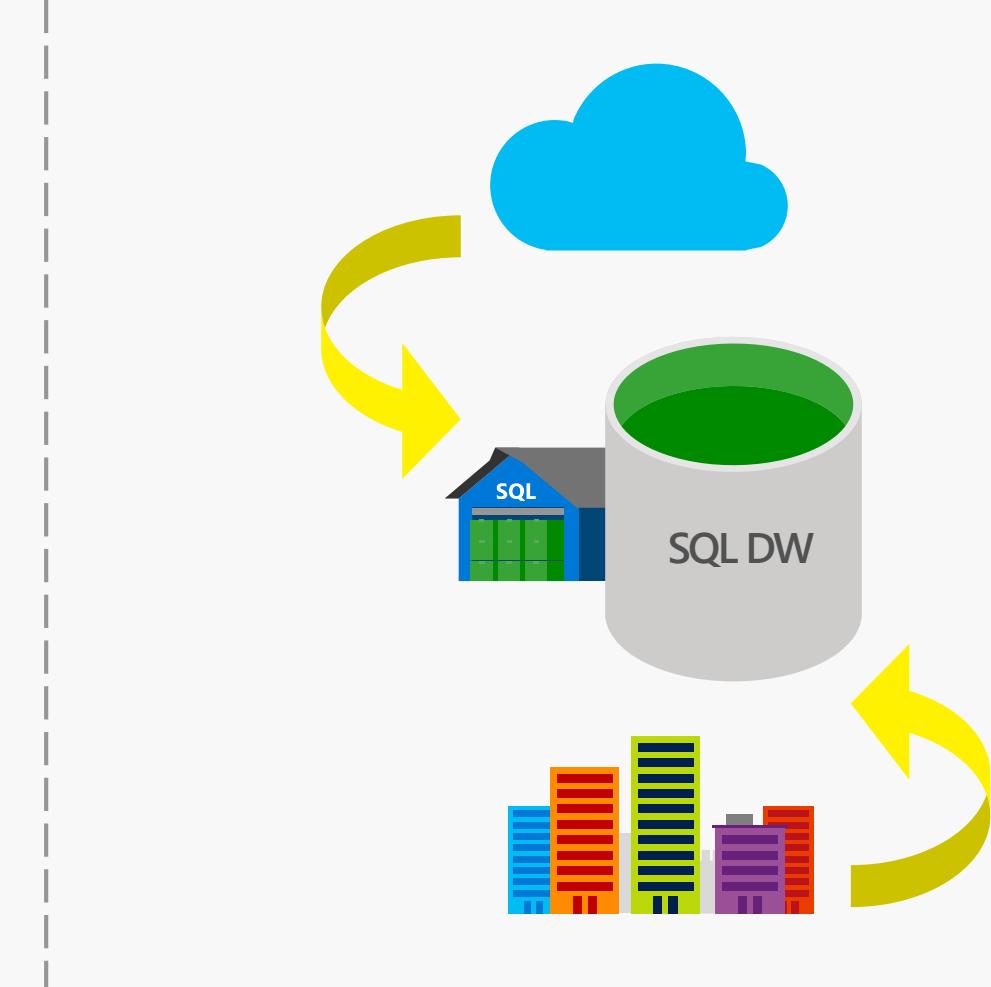
Built on SQL Server
technology using tools
you know love SSMS and
Visual Studio SSDT



Cortana Intelligence Suite



Key differentiators

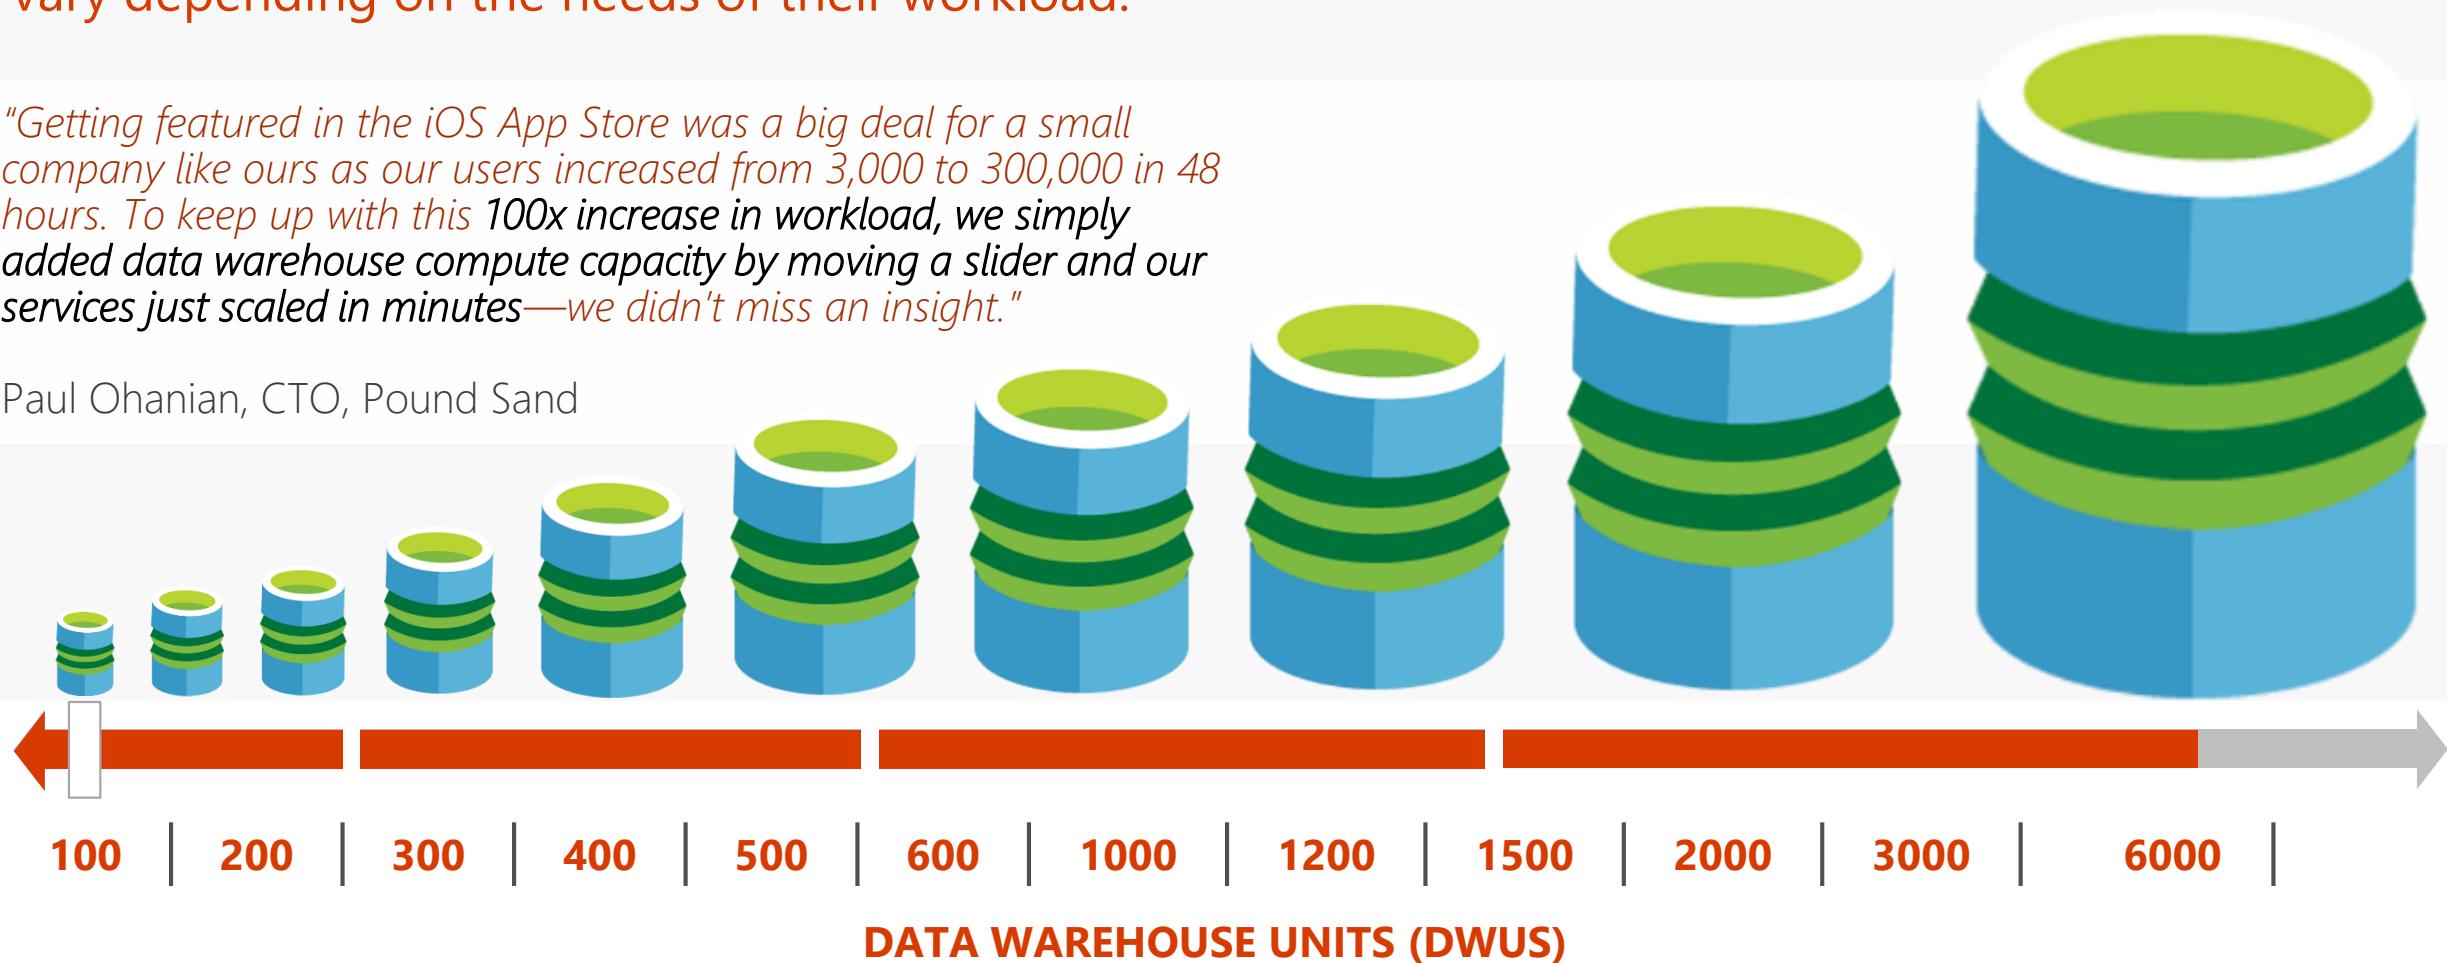


Scale compute on the fly when you need it

Data Warehouse Units (DWUs) are a measure of reserved compute performance or 'power.' A customer's DWU needs can vary depending on the needs of their workload.

"Getting featured in the iOS App Store was a big deal for a small company like ours as our users increased from 3,000 to 300,000 in 48 hours. To keep up with this 100x increase in workload, we simply added data warehouse compute capacity by moving a slider and our services just scaled in minutes—we didn't miss an insight."

Paul Ohanian, CTO, Pound Sand



...pause compute when you don't need it

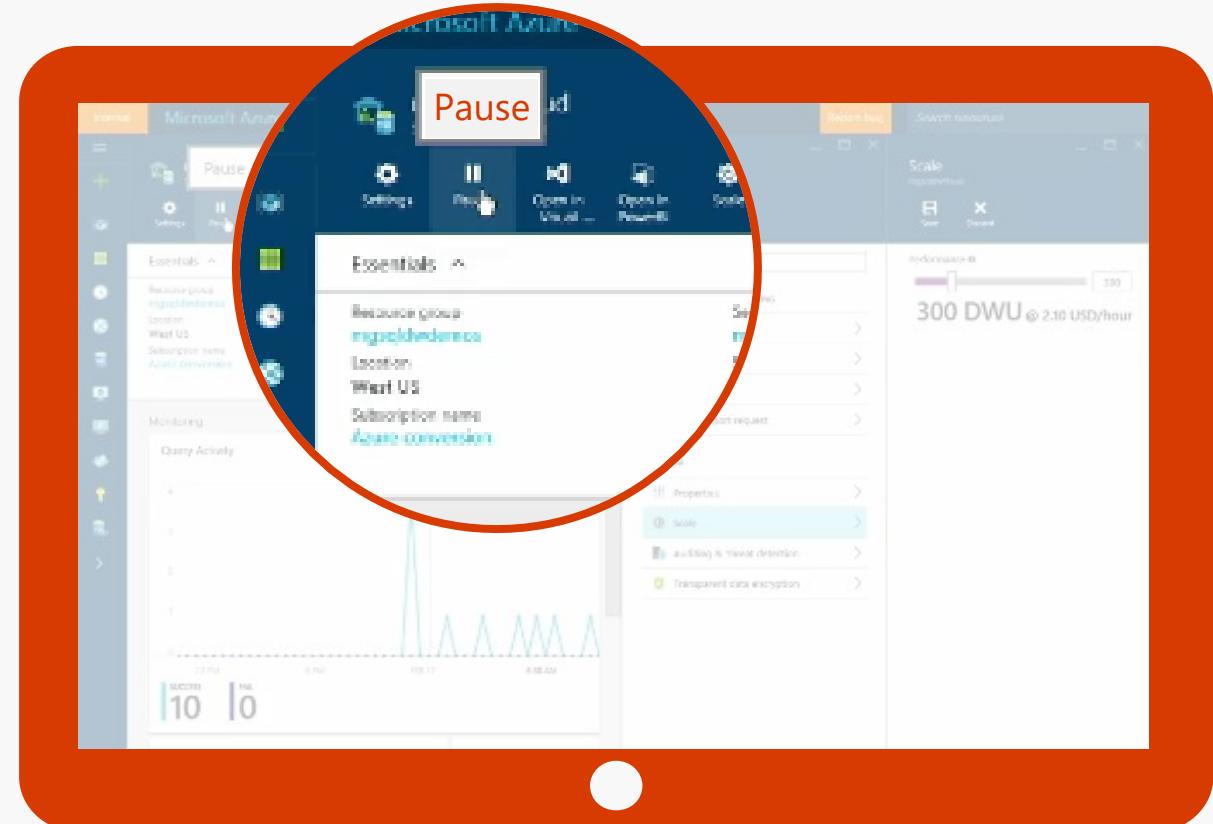
No need to over-provision or over-pay

More freedom to right-size your environment

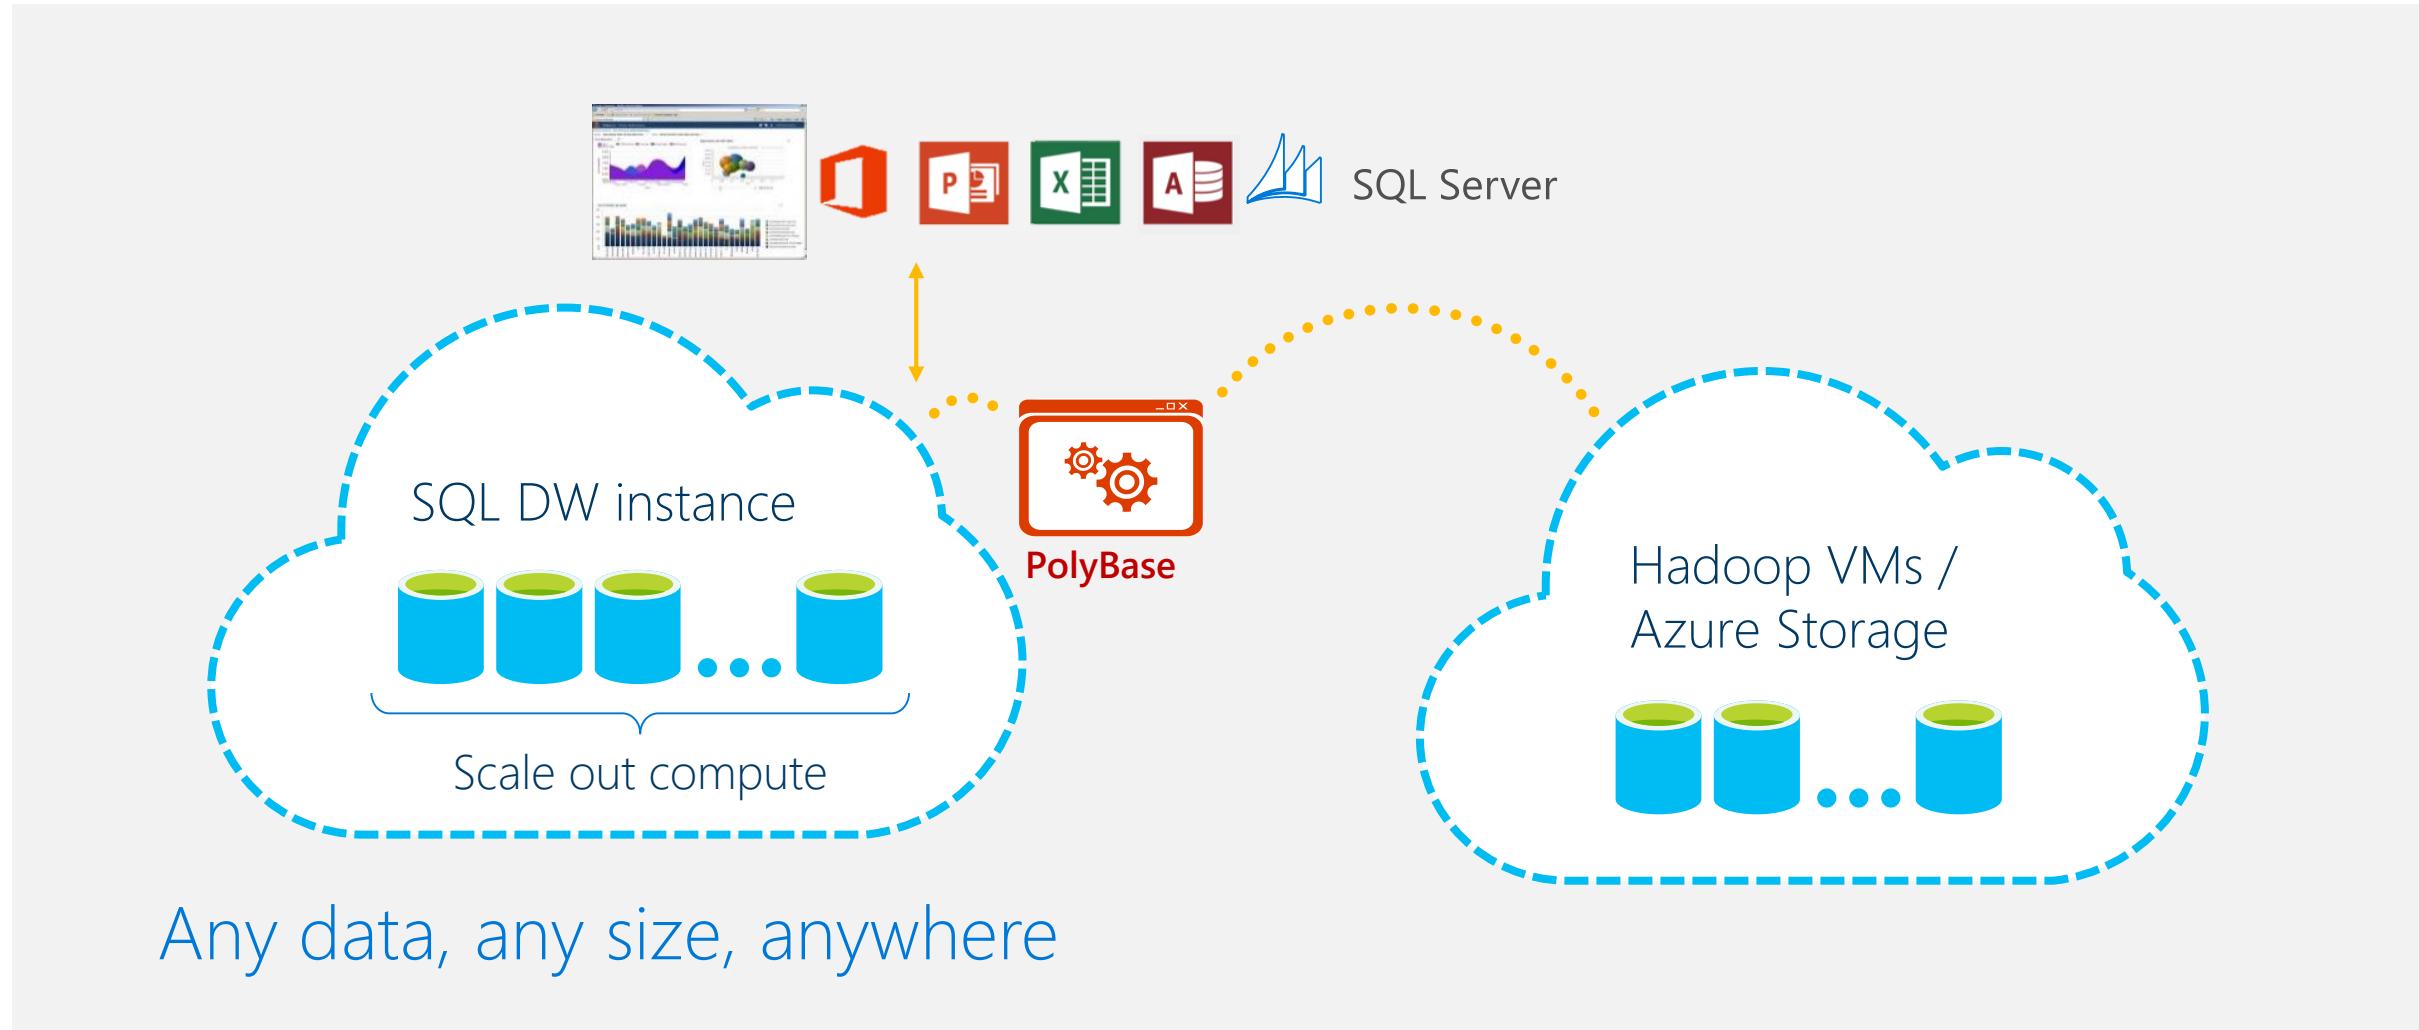
Bring all your data with no tradeoffs on what data to ingest - run compute only on the datasets that matter

"When we learned about the pause and resume capabilities of SQL Data Warehouse, we switched from AWS Redshift, migrating over 25TB of uncompressed data over a week for the simple reason of saving money. To meet our business intelligence requirements, we load data once or twice a month and then build reports for our customers. Not having the data warehouse service running all the time is key for our business and our bottom line."

Bill Sabo, Integral Analytics.



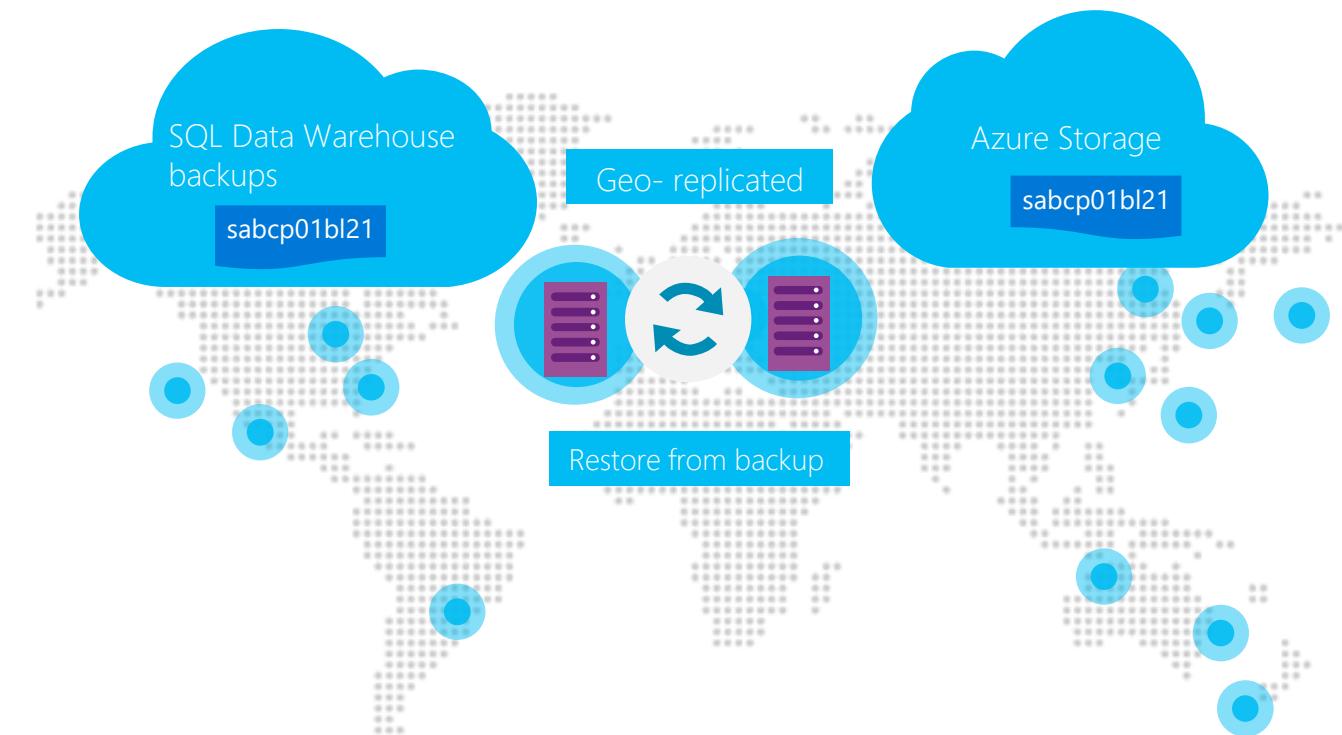
Query unstructured data via PolyBase/T-SQL



Automatic backup and geo-restore

Recover from data deletion or alteration or disaster

- Auto backups, every 4 to 8 hours
- Restore using REST API, PowerShell or Azure Portal
- Local and geo-redundant backups
- Backups retention policy:
 - Auto backup retention for 7 days



Auditing and Threat Detection

Gain real-time insights and streamline compliance-related tasks

RETAIN

an audit trail of selected events and activities

REPORT

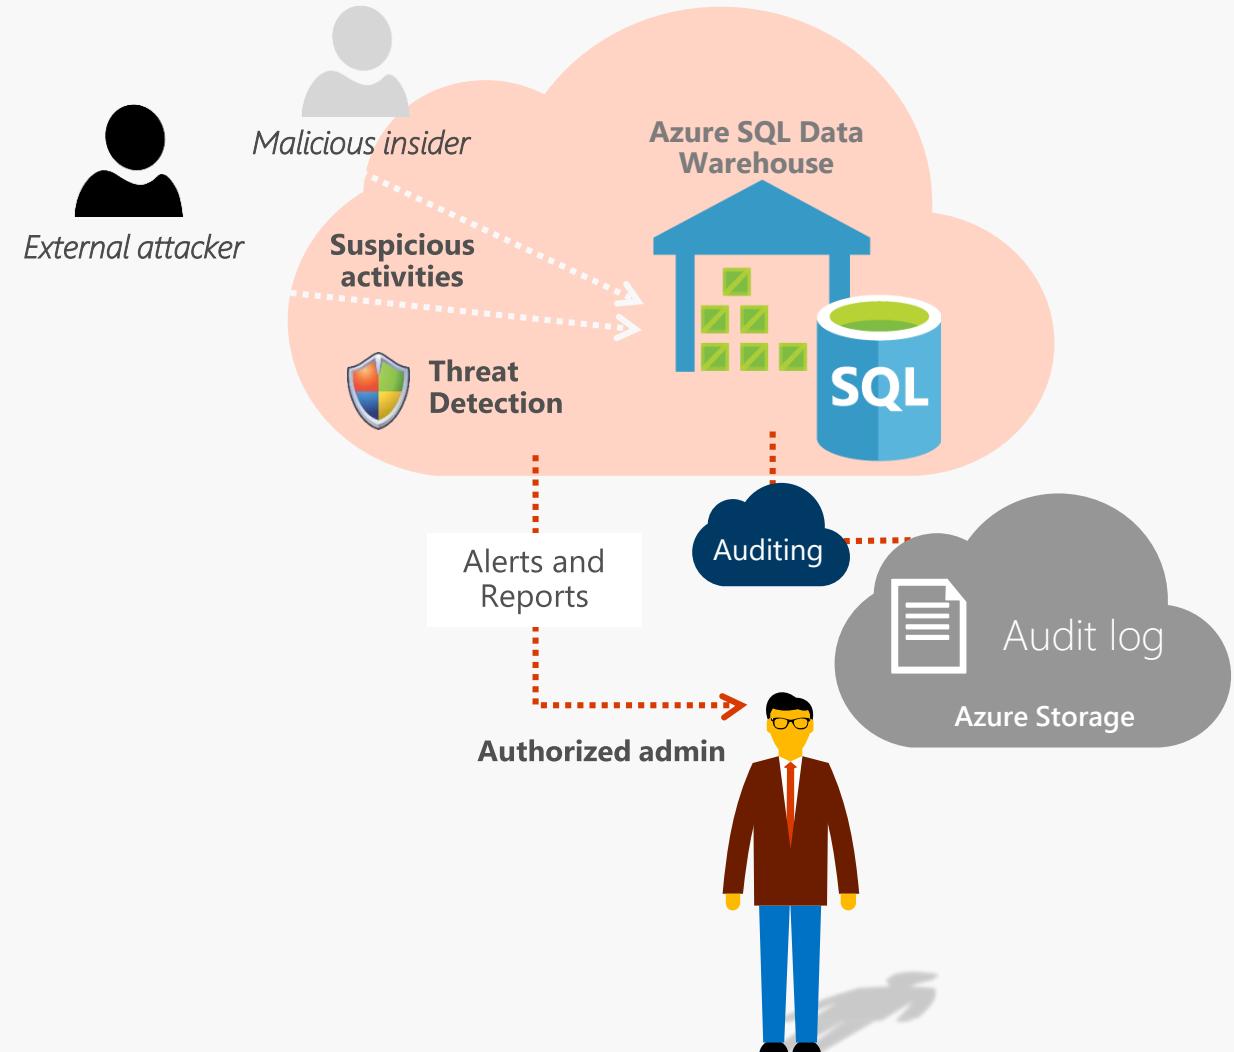
on database activity—preconfigured reports and a dashboard help get you started quickly

ANALYZE

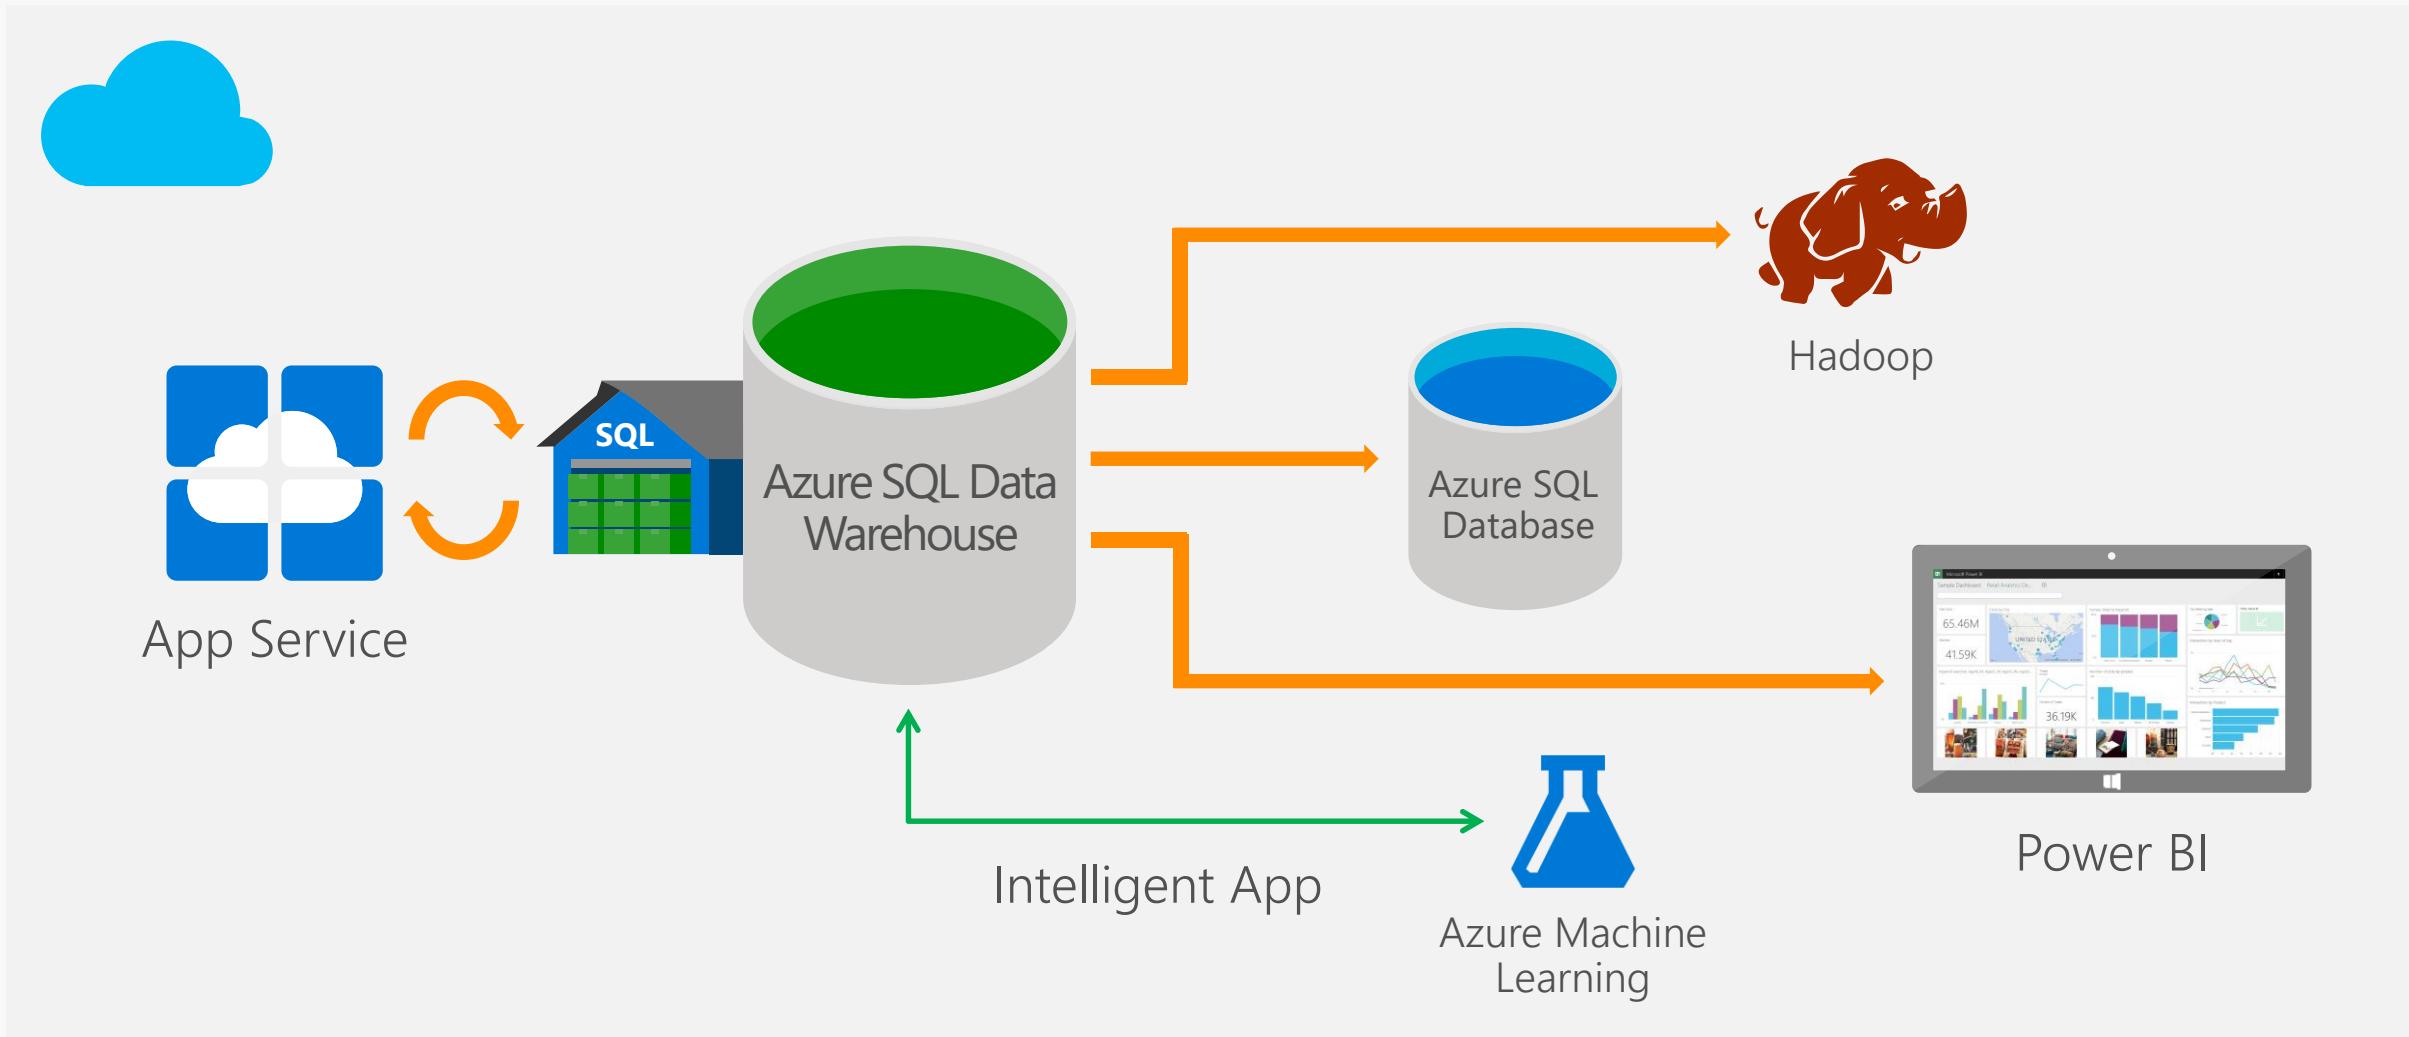
reports to find suspicious events, unusual activities, and trends

RECEIVE PROACTIVE ALERTS

about activities that might indicate potential security threats using the new Threat Detection feature



End to end platform built for the cloud



Your choice of language and tooling

SQL Server		Azure SQL Database			SQL DW		
ADO.NET	JDBC	ODBC	PHP Driver db-lib	Node.js Driver Tedious Node.js Driver	Pymssql FreeTDS	Django Pymssql	Rails TinyTDS
C# VB.NET	Java	C/C++	PHP	JavaScript	Python	Ruby	

SQL DW – Fit for Purpose

- Designed for DW and not OLTP. All the traditional DW workload characteristics apply.
- Not good for singleton DML heavy operations, Example: Clients issuing singleton update, insert, delete.
- Incremental data is loaded regularly by ETL/ELT process in batch mode. Not optimal for real time ingestion.
- Low number of concurrent queries.
- Where data runs in Millions, Billions of rows and TBs in Size.
- Main performance is driven from Massively parallel processing Architecture.

Azure SQL Data Warehouse

Architecture

MPP architecture - Concepts

- **Massively Parallel Processing**
- A divide and conquer, **distributed computing** strategy
- **Shared Nothing** architecture
- Scale-out architecture (versus scale-up)
- Take one big problem & break it up & execute it individually
- Team approach "Many hands make light work"
- Multiple machines (RAM, CPU, Storage) with high bandwidth for single task
- Codeless parallel processing
- Best suited for large data processing
- Out of the box -
 - Task Scheduling, Distribution, Communication & Co-ordination
 - High Availability, Fault Tolerance

MPP architecture - Concepts

Logical layer

- Holds application metadata, Does not persist application data
- Receives intermediate results and performs final aggregation

Physical layer

- Persists application data
- Performs query steps as instructed

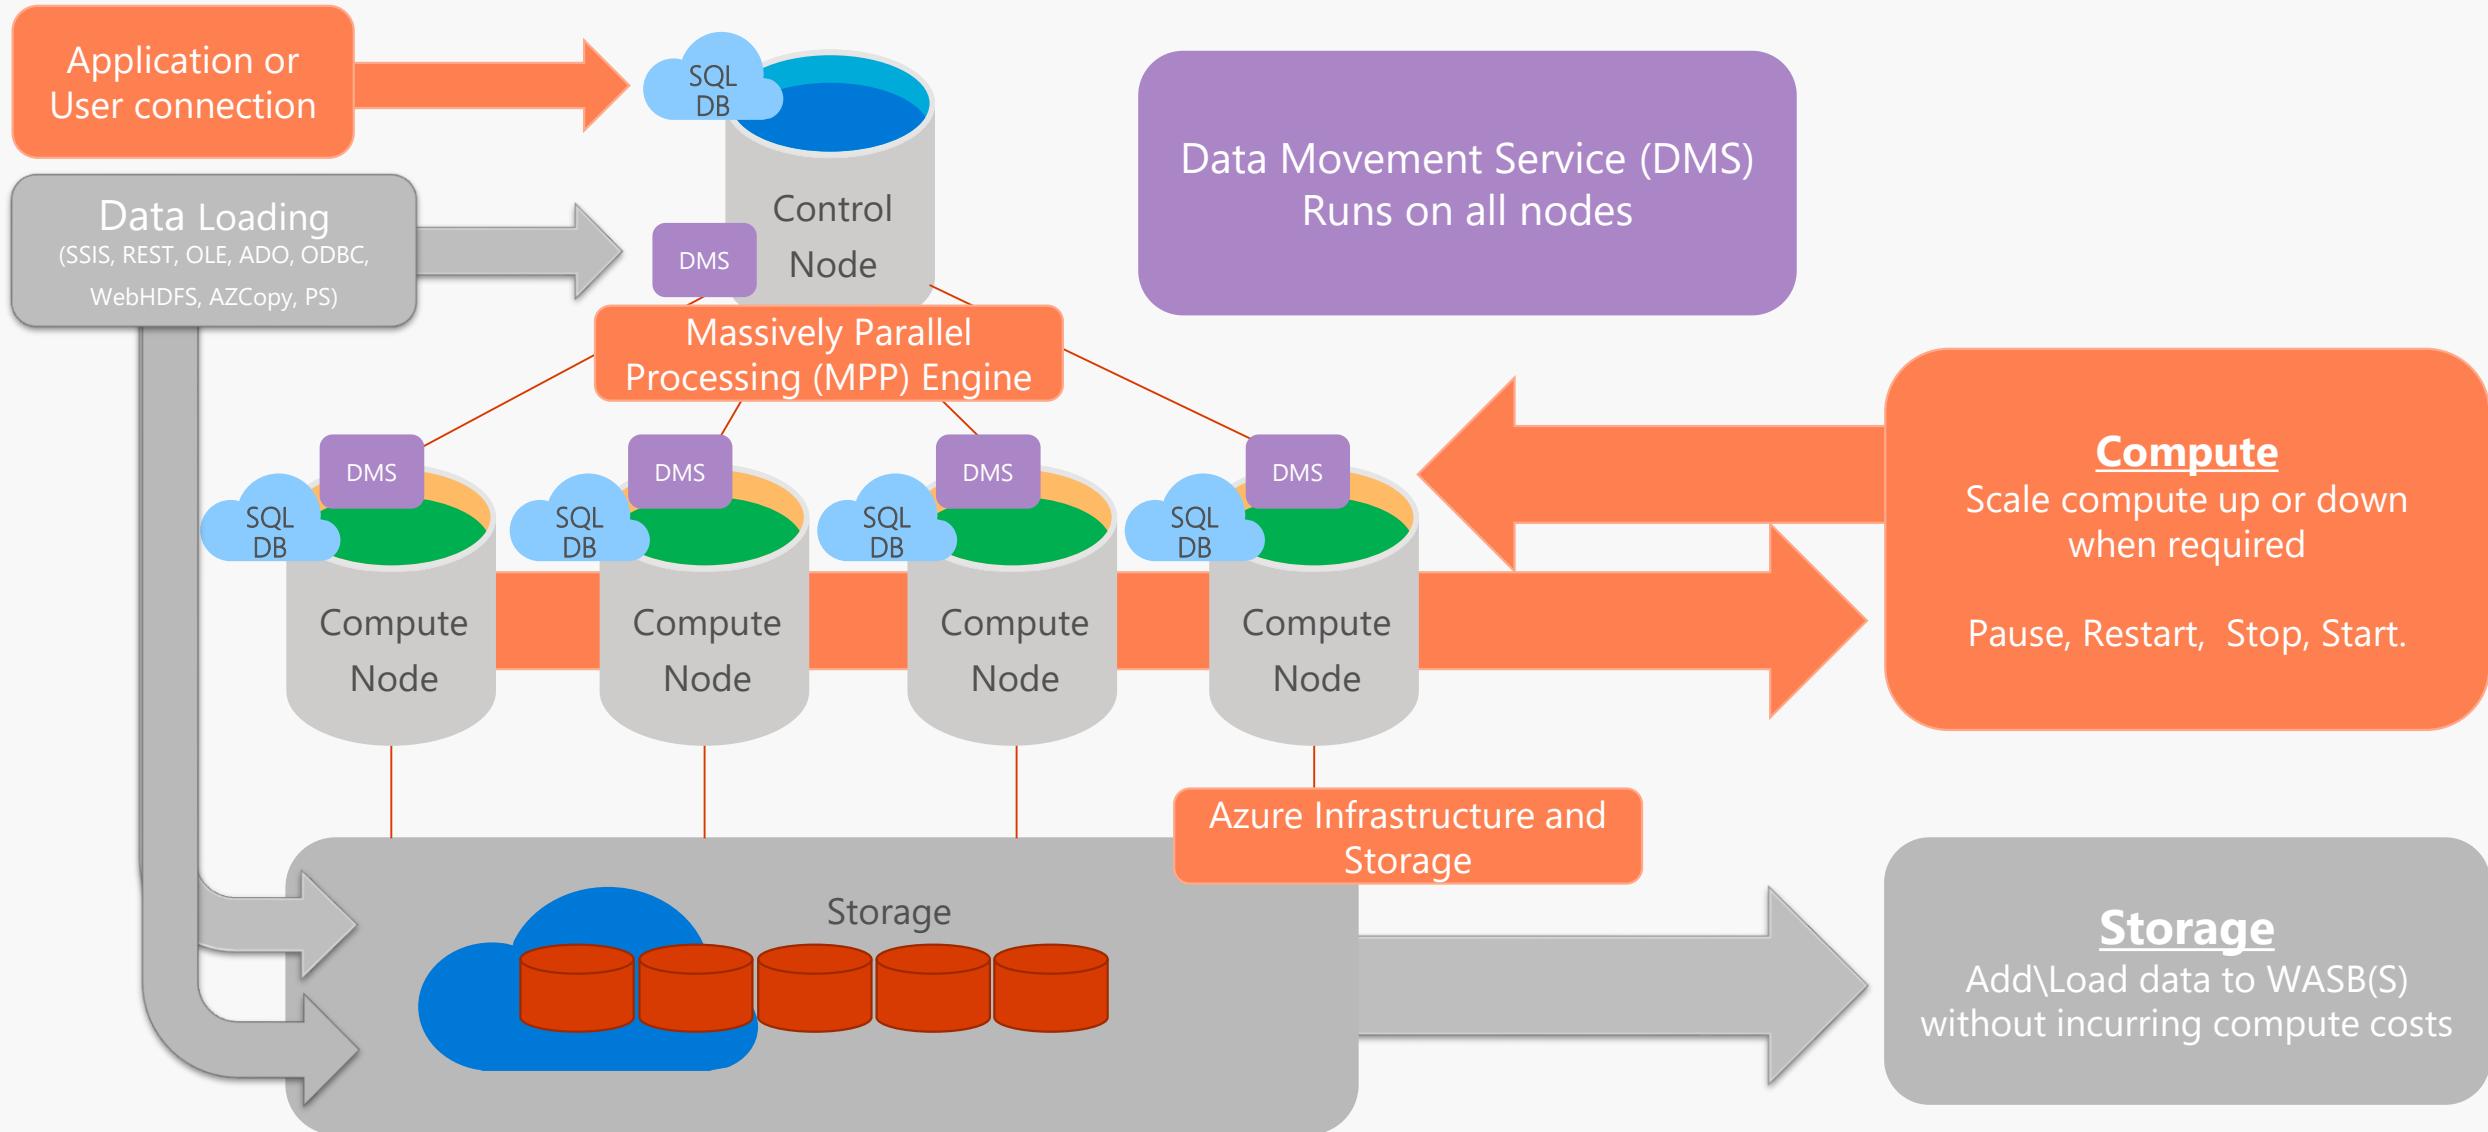
Distributed Query engine

- Parallel processing coordination

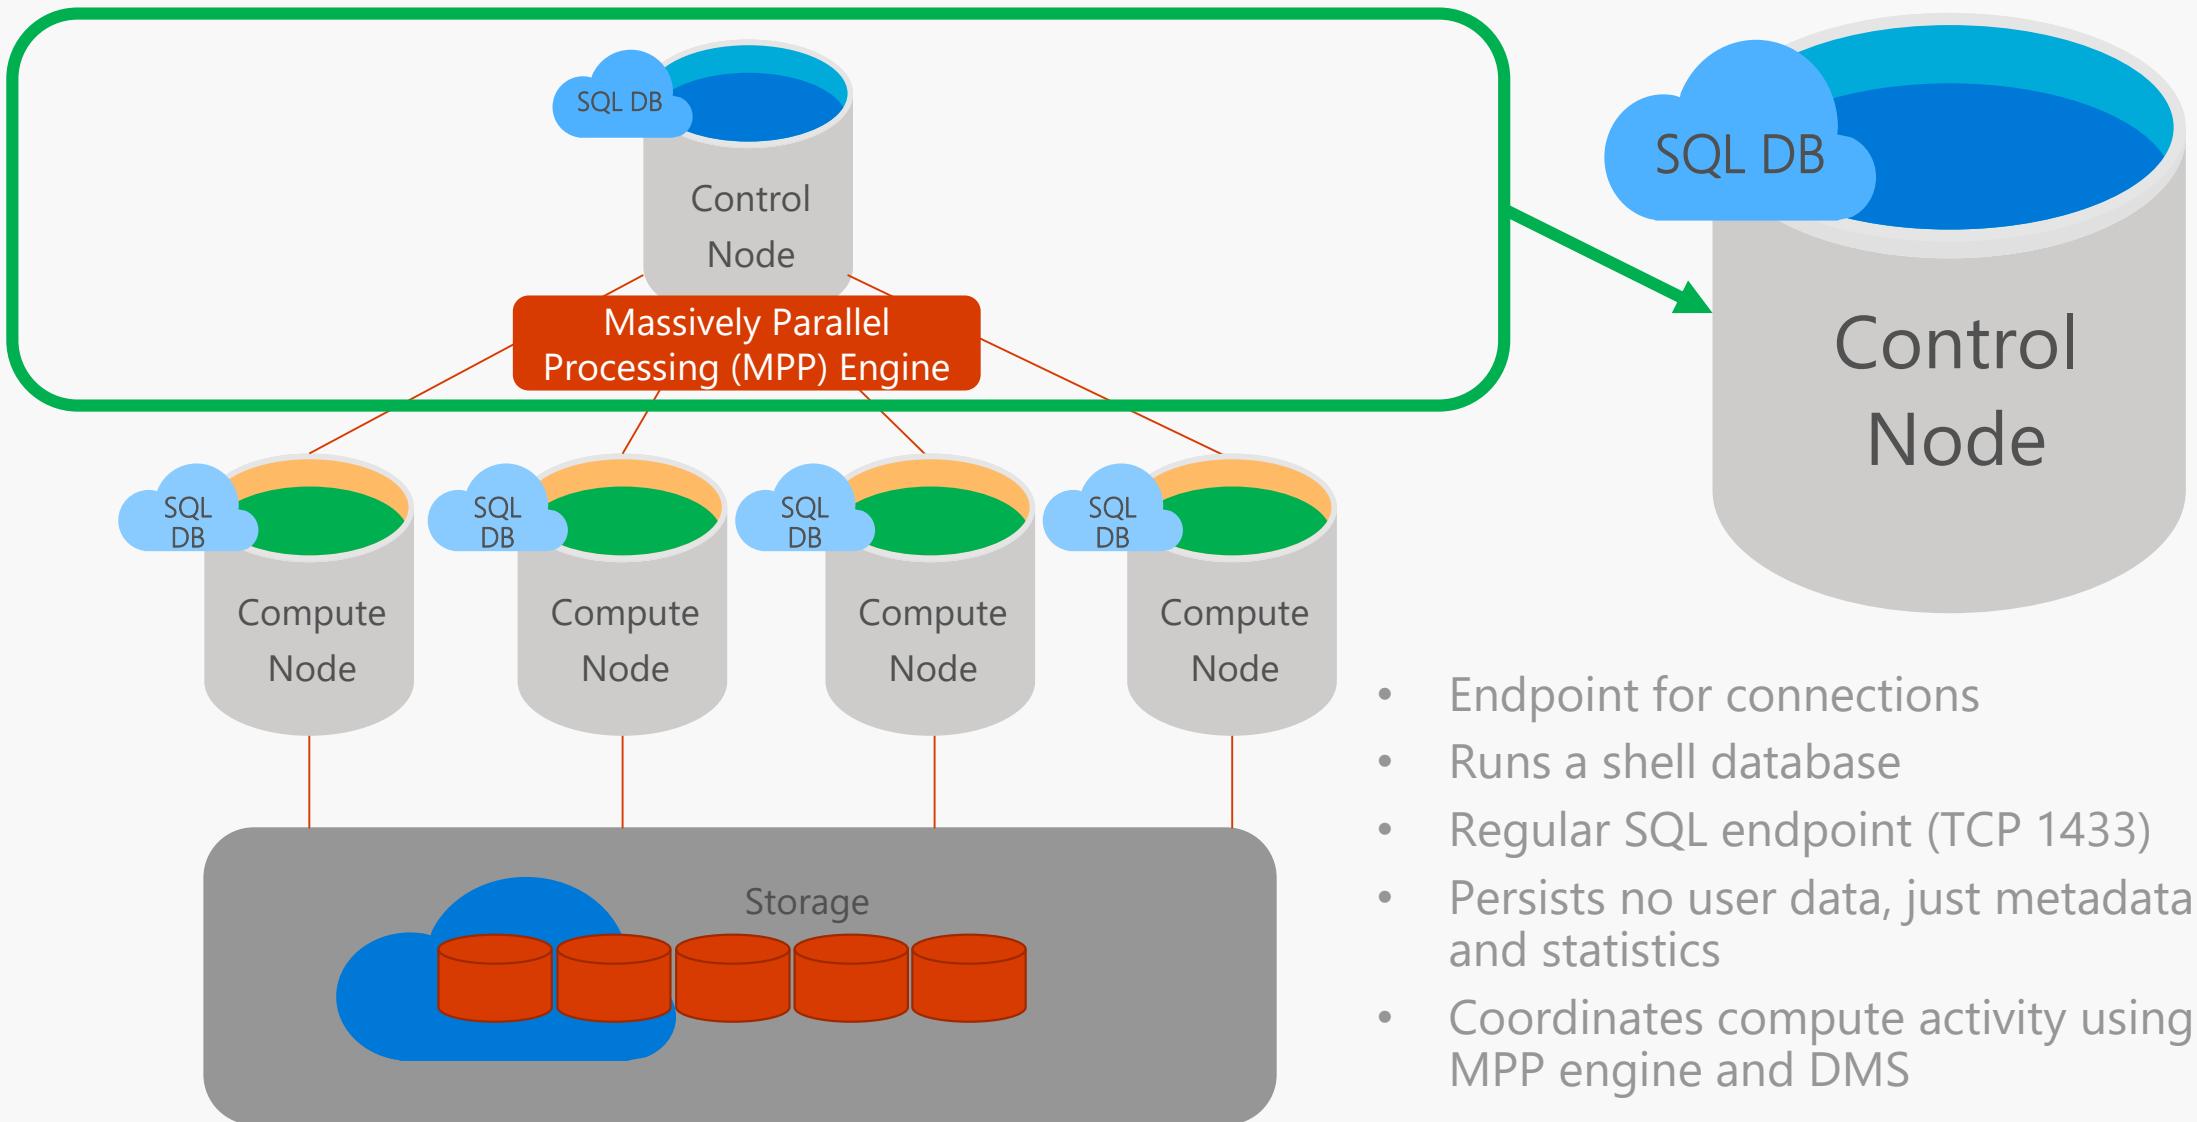
Lots of machines!

- Lots of small databases running in parallel

Azure SQL Data Warehouse Architecture



Azure SQL Data Warehouse – Control Node



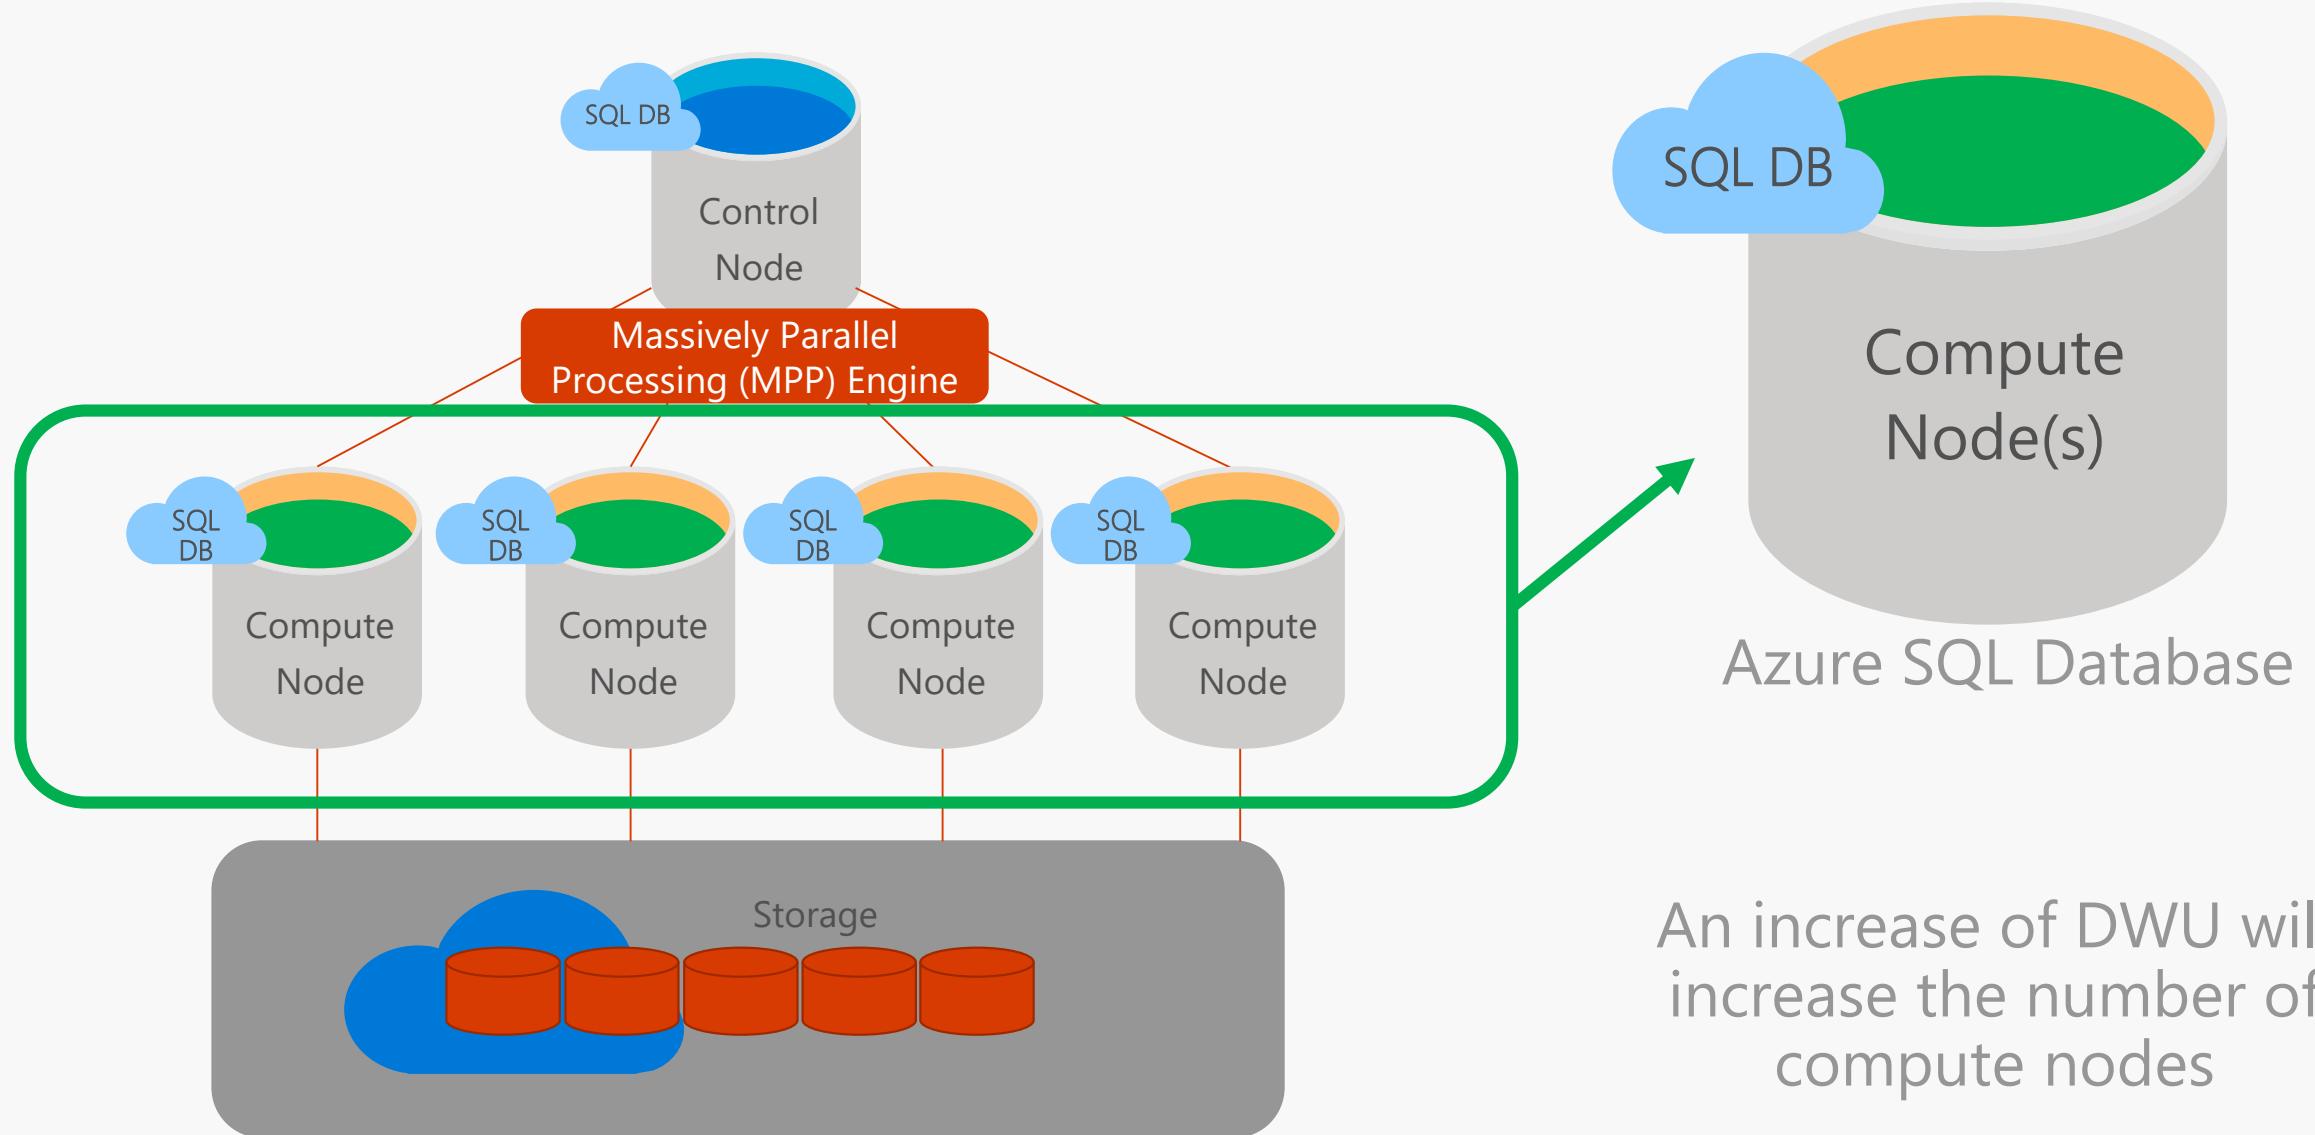
Control node

- Accepts requests from user and transforms it into separate queries that run on each compute node in parallel
- Coordinates all of the data movement and computation required to run parallel queries
- Interprets requests for scale-out
- Orchestrates actions / steps
- Final computation
- Returns result

- ✓ Runs SQL Server
- ✓ Stores metadata, statistics
- ✓ Performs final computation

Requires distributed query engine

Azure SQL Data Warehouse - Compute Nodes



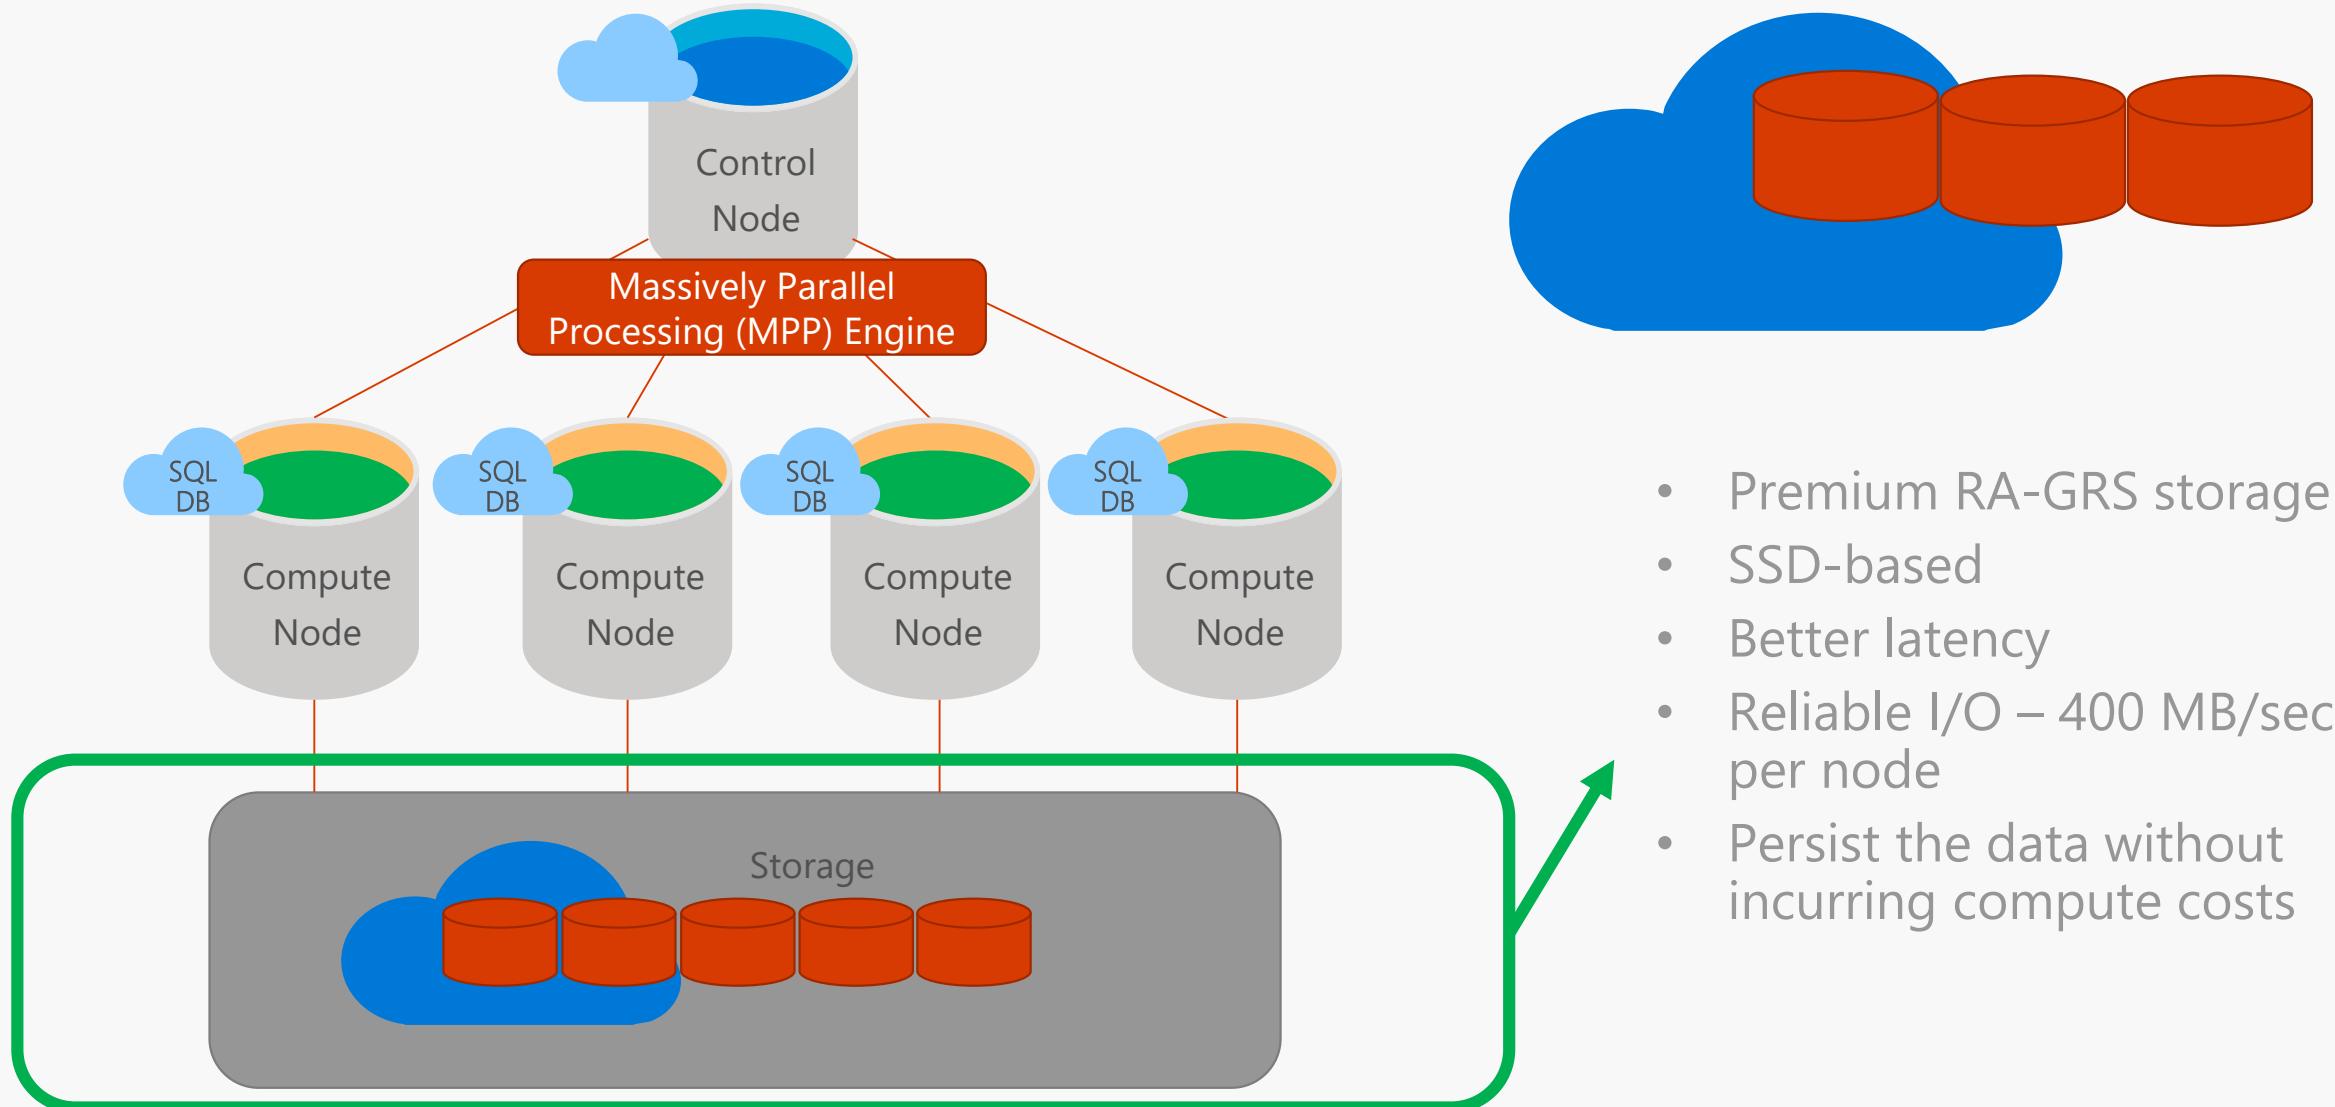
Compute node

- Power behind the SQL DW; Performs the heavy lift
 - Accepts requests from control node
- Is a highly tuned SMP System
 - Optimises requests from control node
- SQL Databases that store data
- Workers that run parallel queries on data

No direct user interaction

Uses SQL Server

Azure SQL Data Warehouse – premium storage

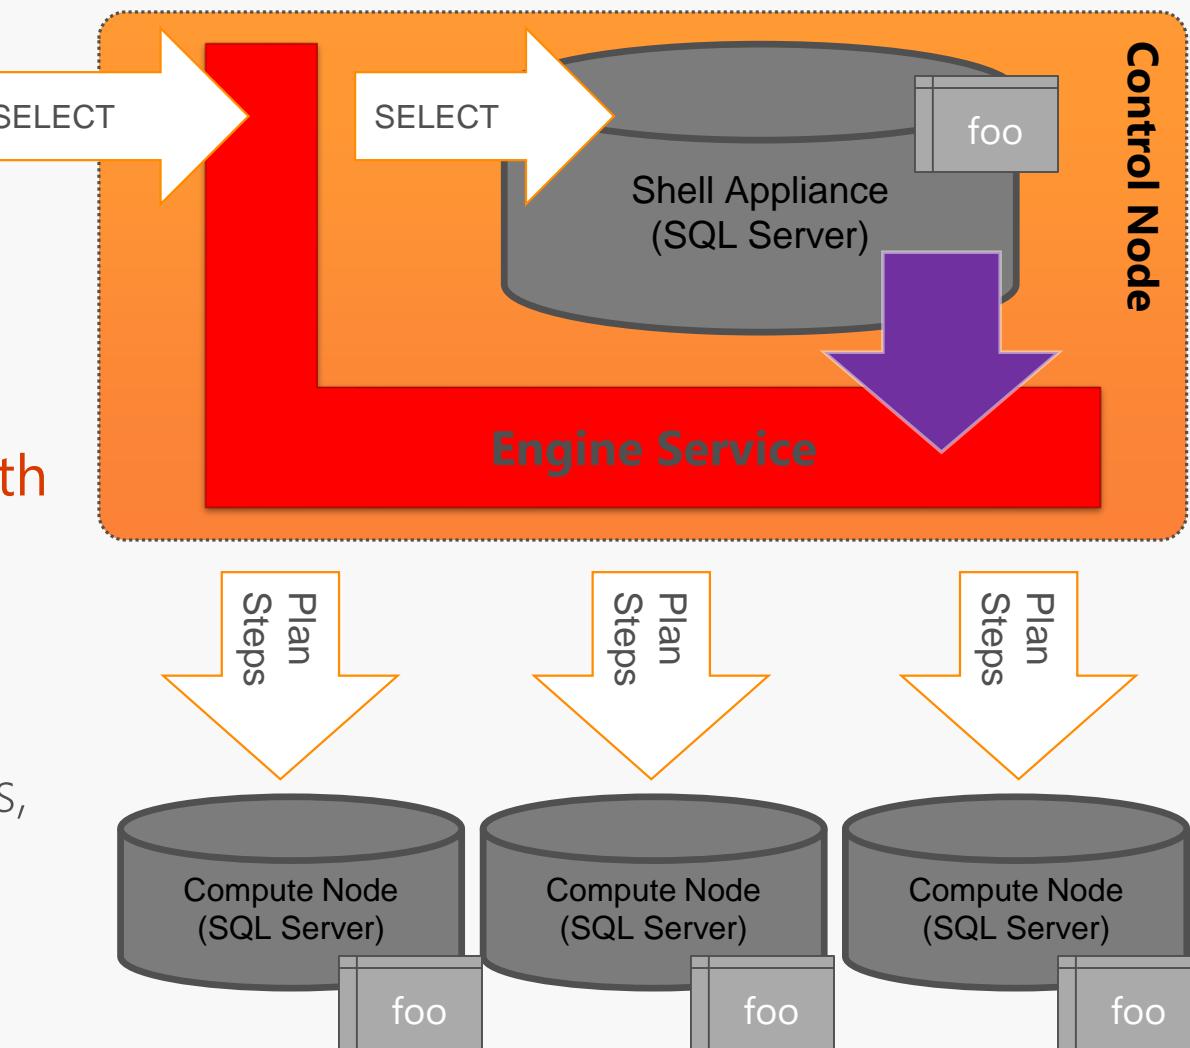


SQL DW Storage - Page blobs

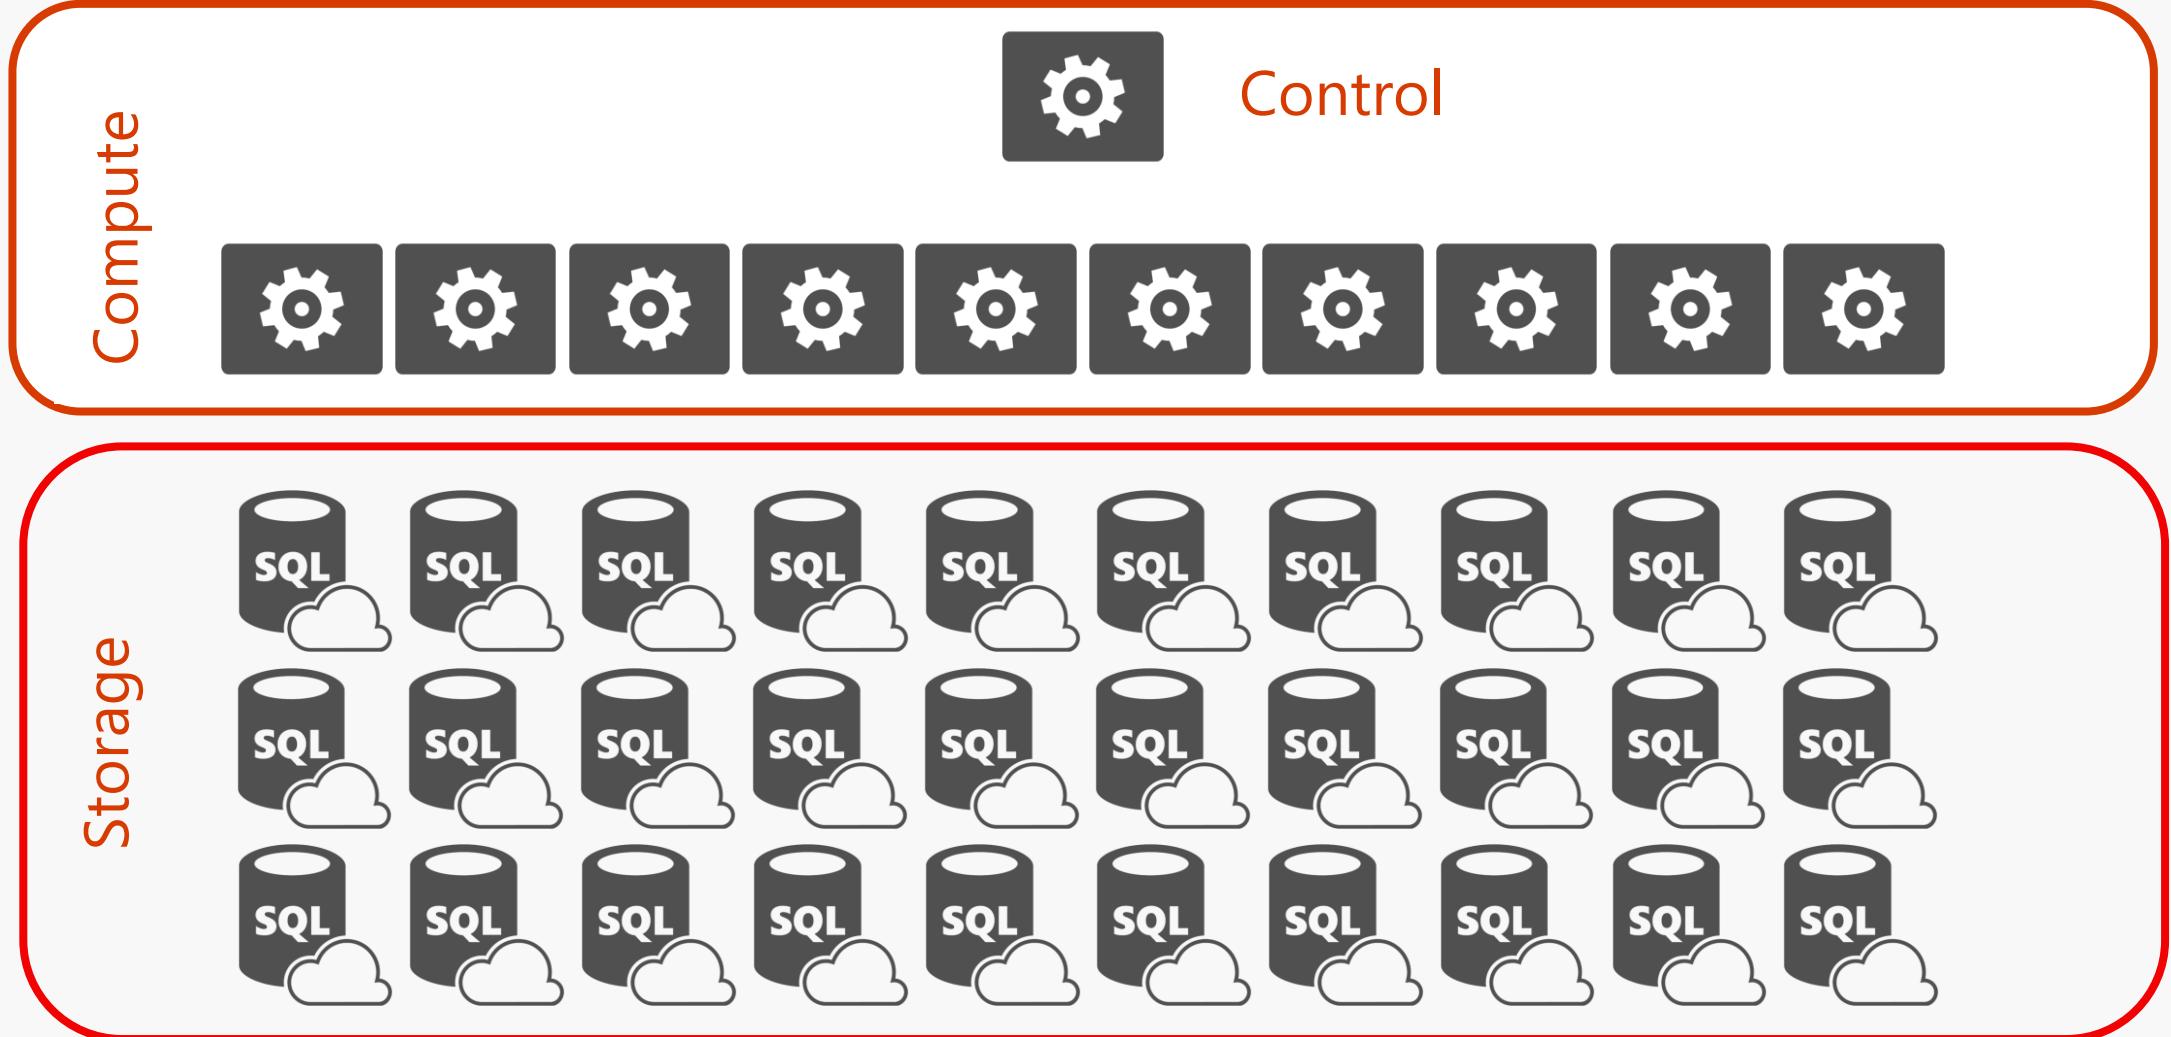
- Azure SQL DW is backed by SSD based page blobs
- Page blobs – optimized for representing IaaS disks
- Supports random writes, up to 1 TB in size
- Premium page blobs – optimized for I/O intensive workloads
- Premium page blob is LRS (Locally Redundant Storage) only
- *<https://docs.microsoft.com/en-us/rest/api/storageservices/fileservices/understanding-block-blobs--append-blobs--and-page-blobs>*

SQL DW Control Architecture

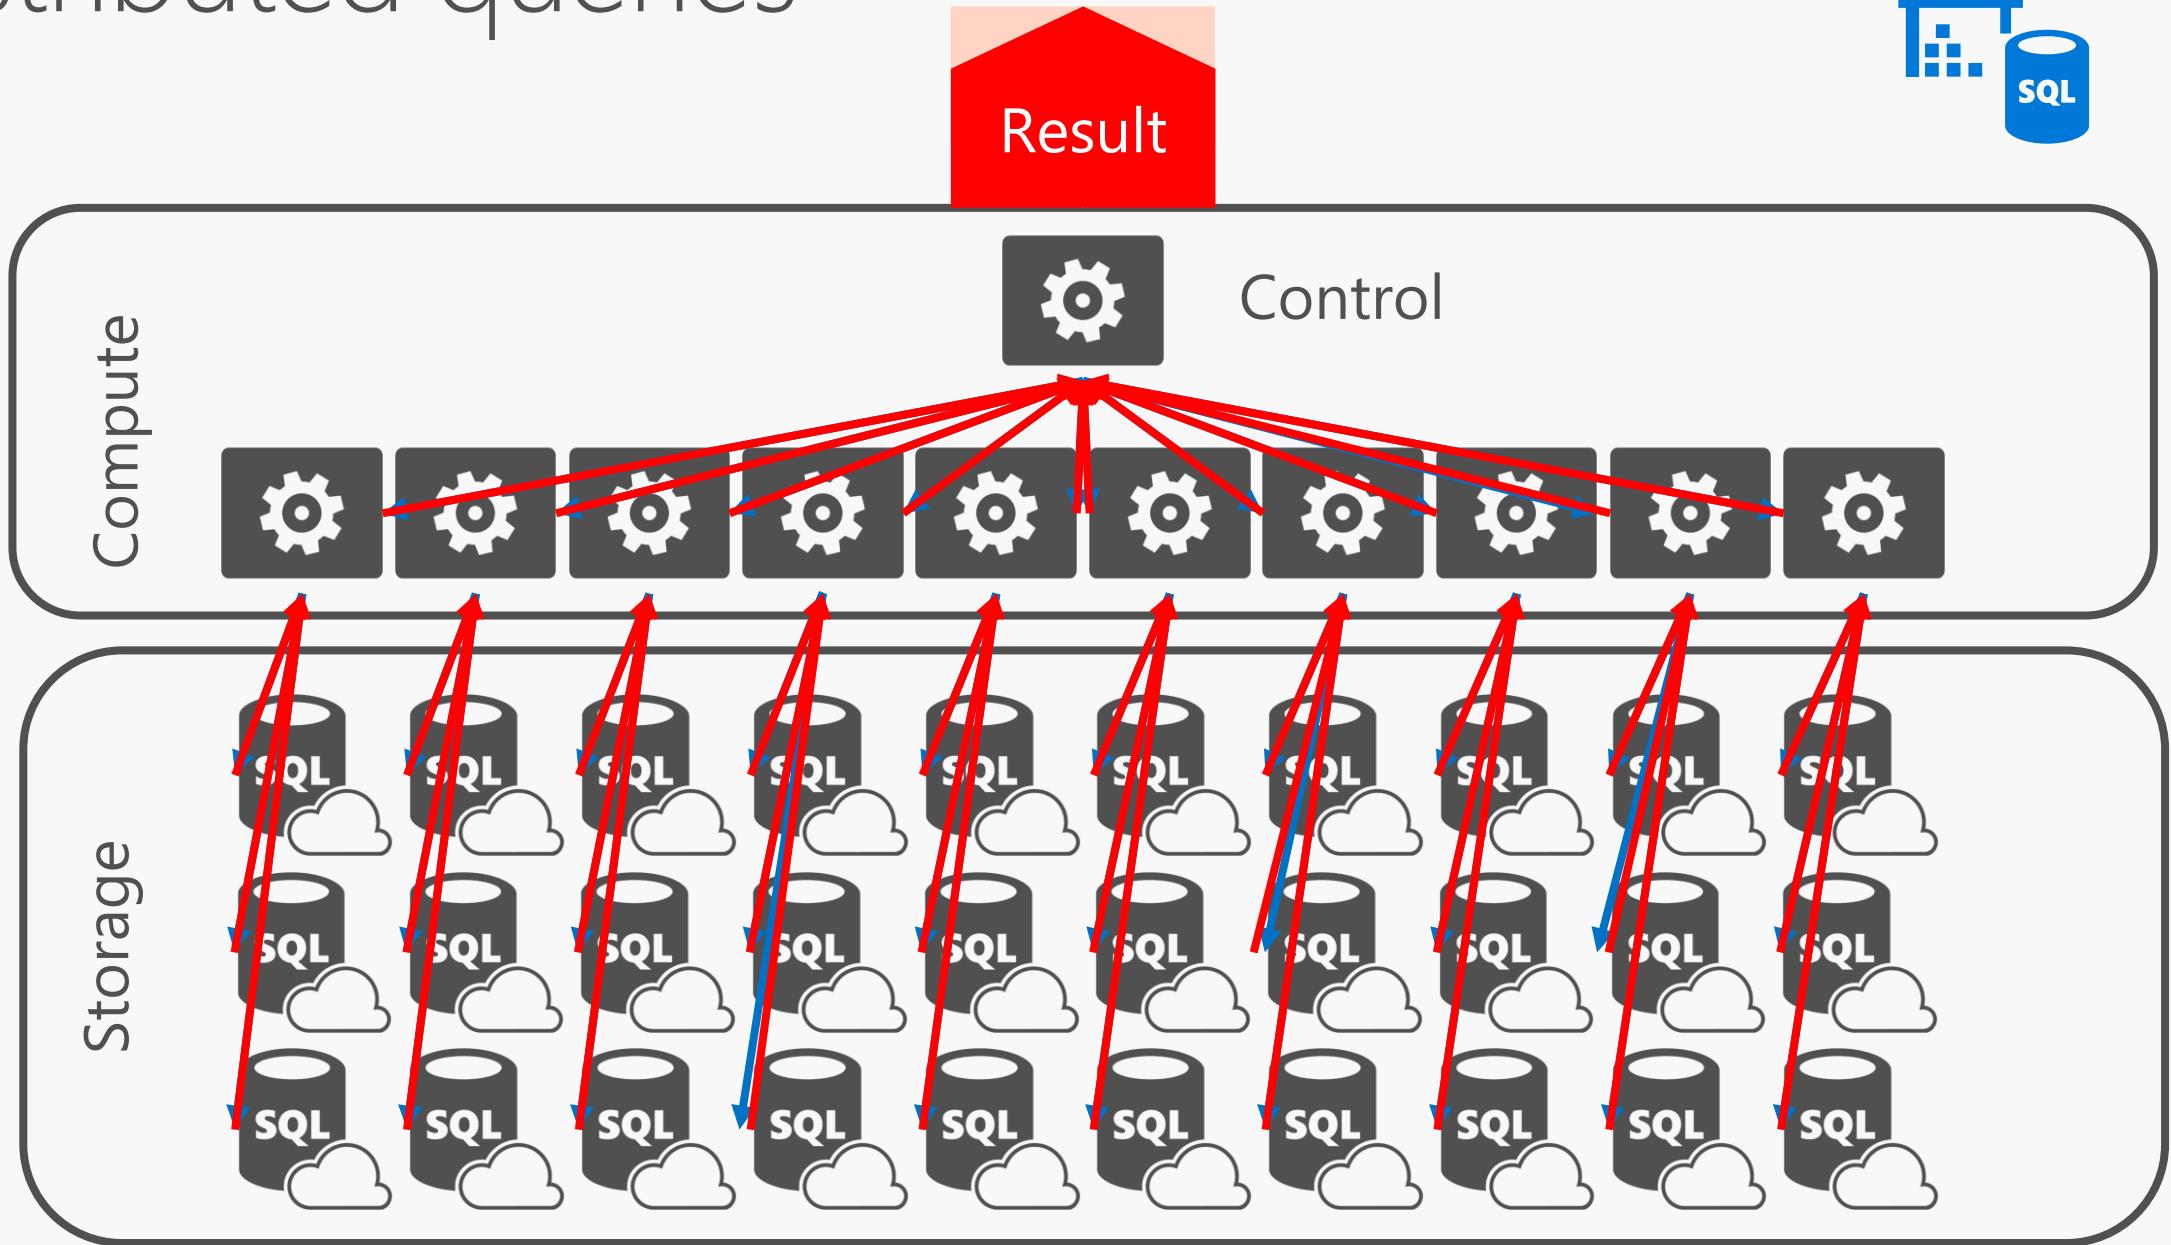
- SQL DW uses SQL Server on the Control node to run a “shell appliance”
- Every database with all its objects exists in the shell appliance as an **empty** “shell,” lacking the user data (which sits on all the compute nodes)
- Every DDL operation is executed against **both** the **shell and the compute nodes**
- Large parts of **basic RDBMS functionality** provided by that shell
 - Authentication and authorization of queries, but also the full security system
 - Schema binding
 - Metadata catalog



Logical overview



Distributed queries



Simple Example

```
SELECT COUNT_BIG(*)  
FROM dbo.[FactInternetSales];
```



```
SELECT SUM(*)  
FROM dbo.[FactInternetSales];
```



Control

Compute

```
SELECT COUNT_BIG(*)  
FROM dbo.[FactInternetSales];
```



```
SELECT COUNT_BIG(*)  
FROM dbo.[FactInternetSales];
```



```
SELECT COUNT_BIG(*)  
FROM dbo.[FactInternetSales];
```



```
SELECT COUNT_BIG(*)  
FROM dbo.[FactInternetSales];
```



Partial Parallelism

- Performed in parallel across compute nodes
- Performed in series across the distributions
- Guarantees transactional behaviour
- Examples: INSERT, UPDATE, DELETE

Full Parallelism

- Performed in parallel across compute nodes
- Performed in parallel across the distributions
- Used for
 - New object creation
 - Reading data from a distributed table
- Examples:
 - CTAS (Create Table As Select)
 - SELECT Statements

Azure SQL Data Warehouse

Scaling

Data warehouse unit (DWU)

Measure of power

Simply buy the query performance you need, not just hardware

Transparency

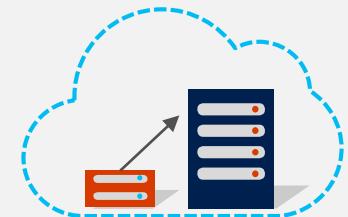
Quantified by workload objectives: how fast rows are scanned, loaded, copied

On demand

First DW service to offer compute power on demand, independent of storage

100 DWU*	Scan Rate	3.36M row/sec
	Loading Rate	130K row/sec
	Table Copy Rate	350K row/sec

Scan 1B rows*	100 DWU	=	297 sec
	400 DWU	=	74 sec
	800 DWU	=	37 sec
	1,600 DWU	=	19 sec

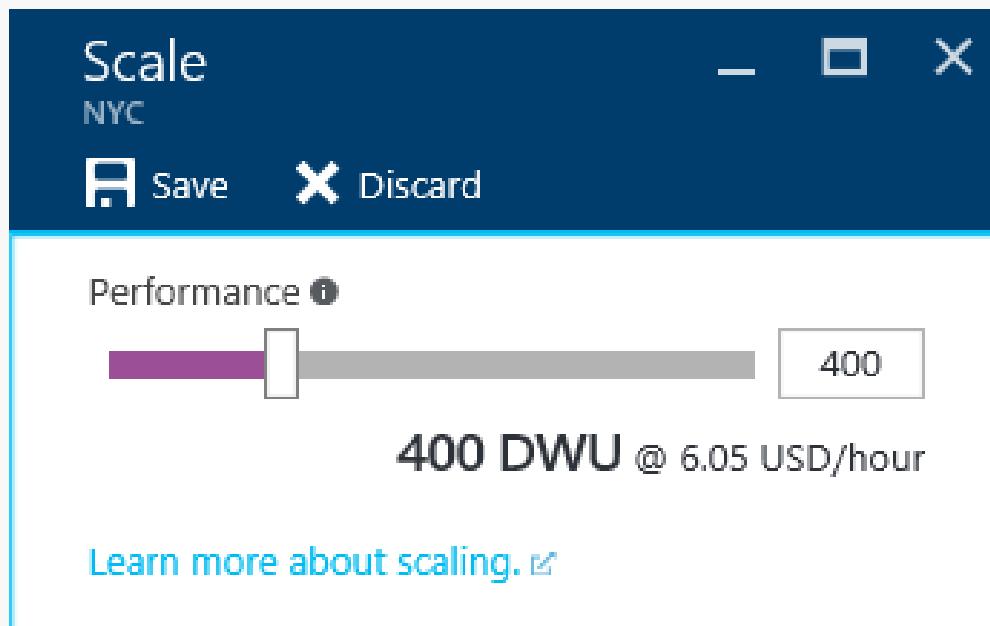


* Preliminary estimates; actual results may change

Data Warehouse Units (DWU)

DWU
DW100
DW200
DW300
DW400
DW500
DW600
DW1000
DW1200
DW1500
DW2000
DW3000
DW6000

```
ALTER DATABASE ContosoDW MODIFY  
(service_objective = 'DW1000');
```



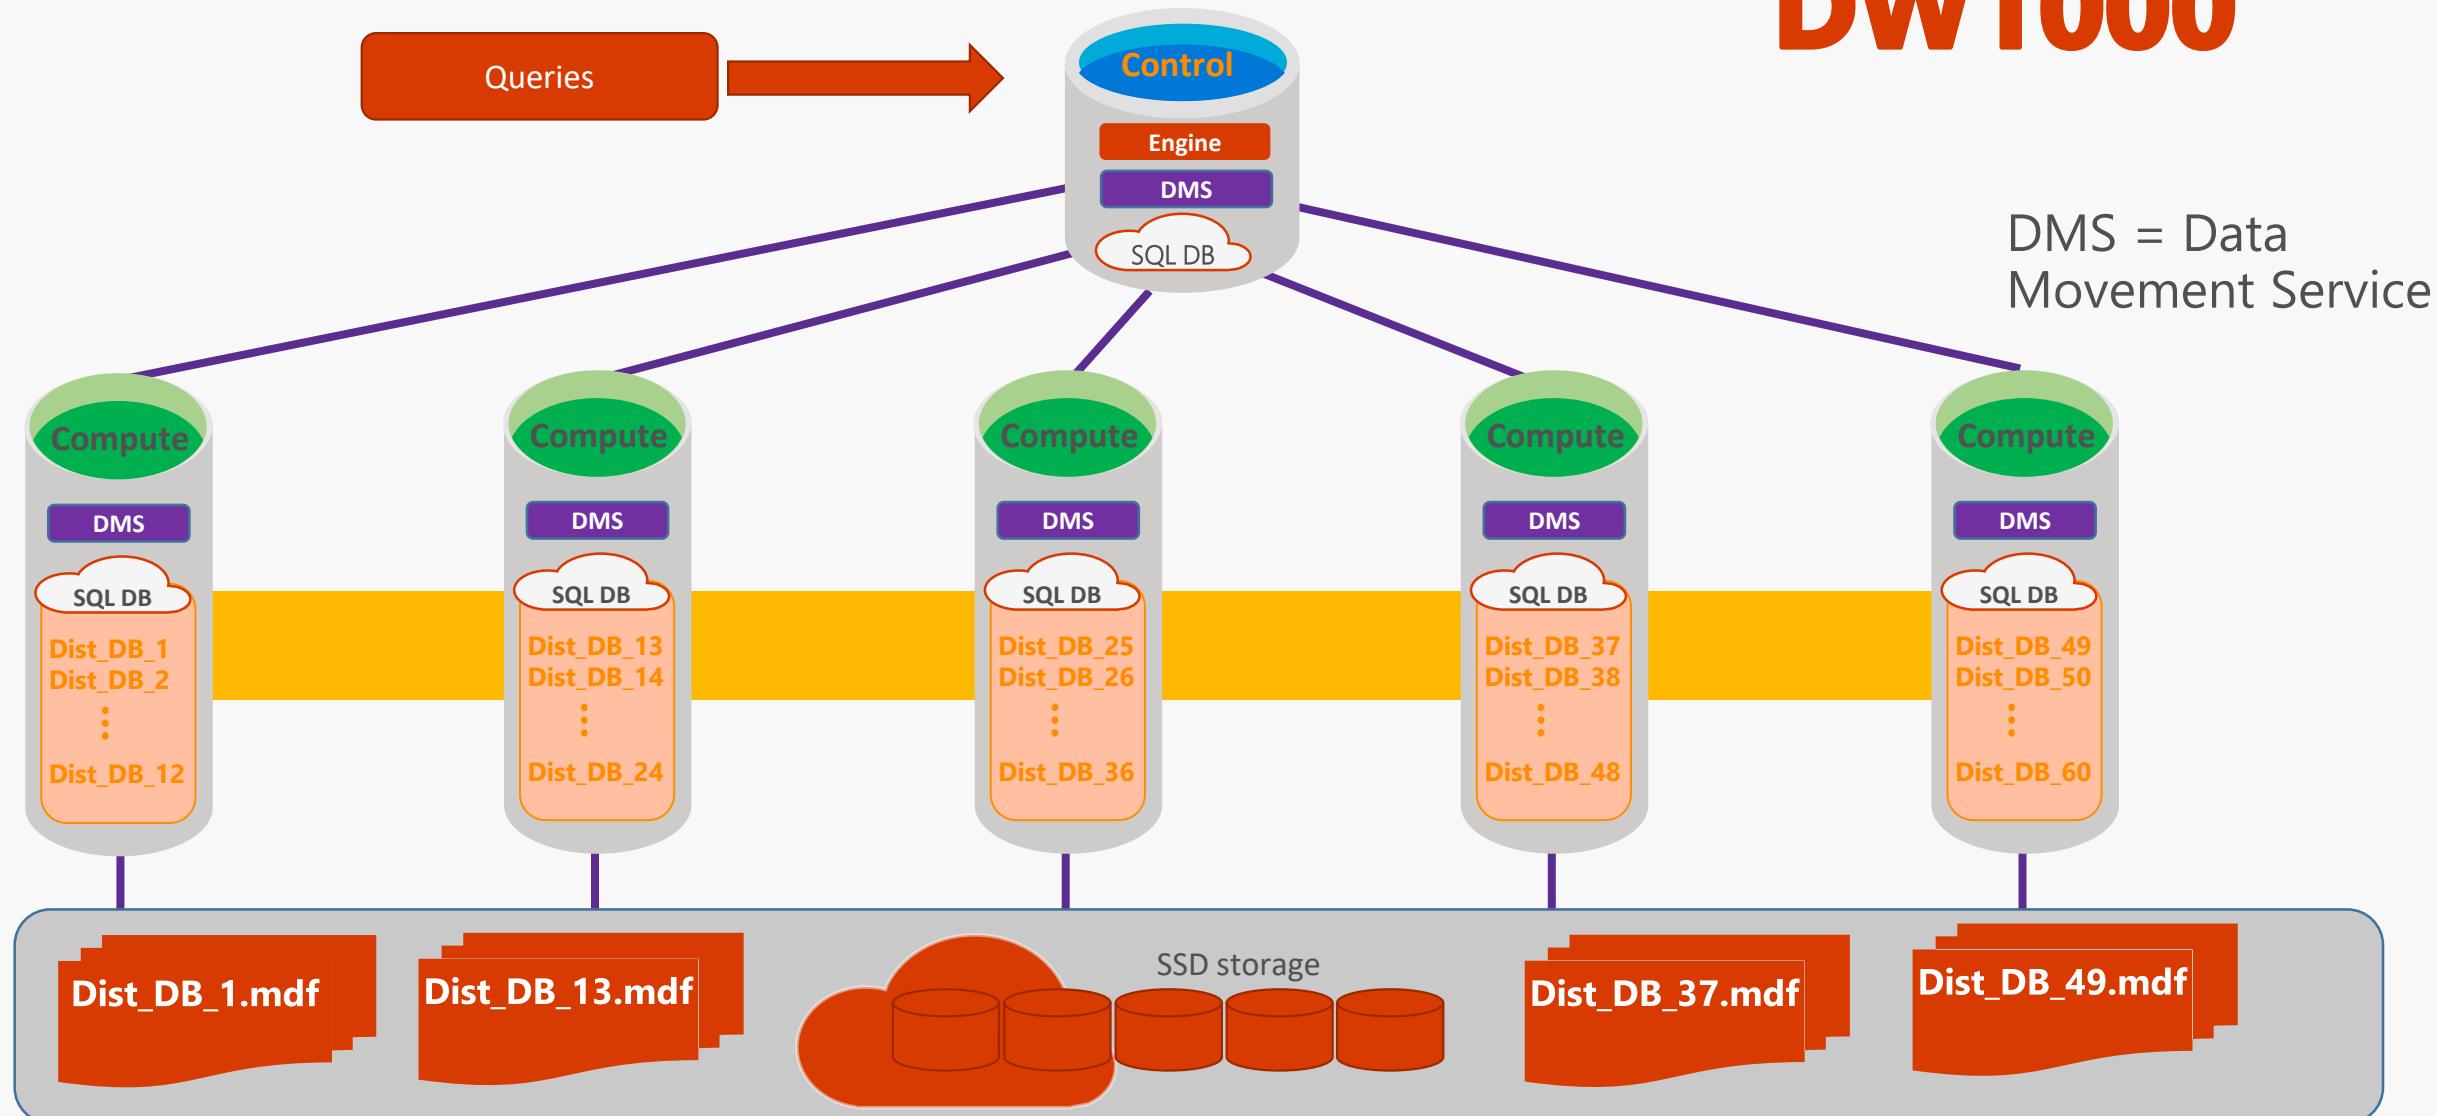
CPU

RAM

I/O

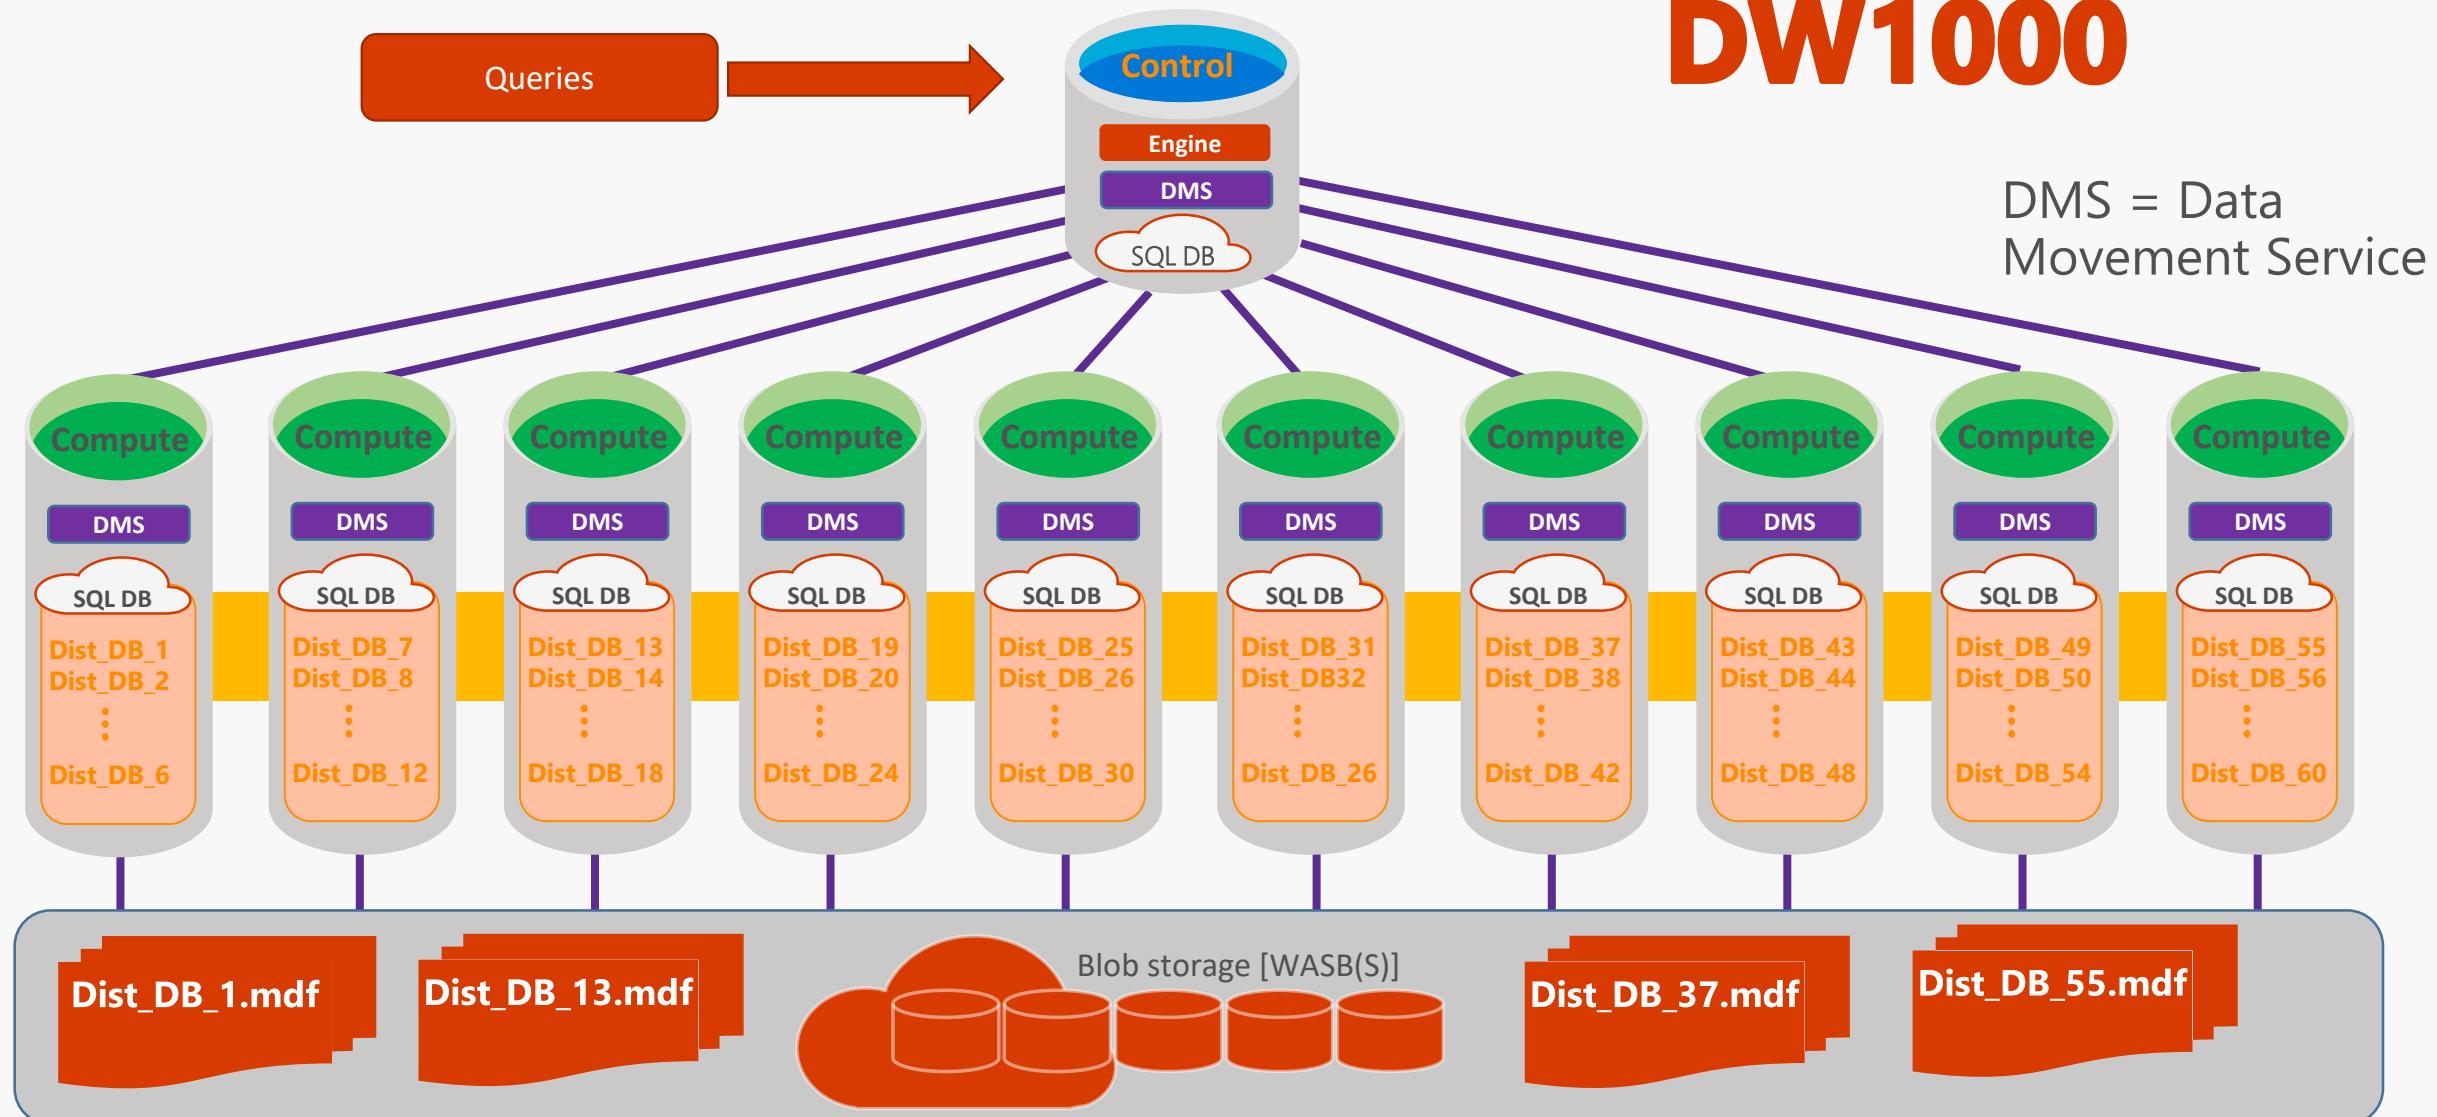
SQL DW Scaling

Changing to
DW500 to
DW1000



SQL DW Scaling

Challenging
DW1000



Azure SQL Data Warehouse

Sizing

Compute Sizing

Sizing by capacity

Recommended starting point



Flexibility to select any range of DWUs



>160 TBs



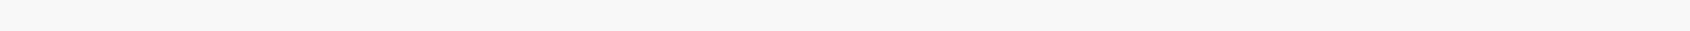
80-160 TBs



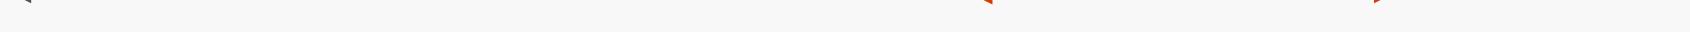
60-80 TBs



48-60 TBs



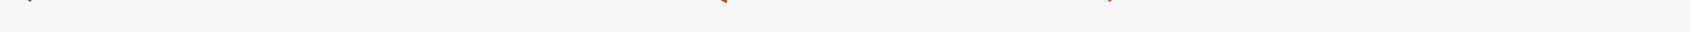
36-48 TBs



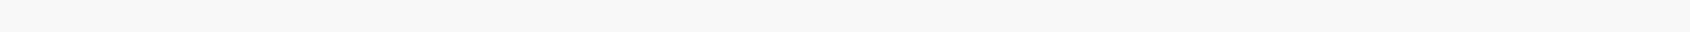
20-36 TBs



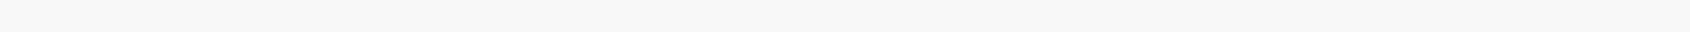
16-20 TBs



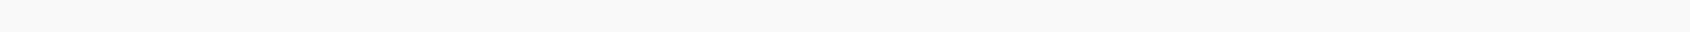
12-16 TBs



8-12 TBs



4-8 TBs



0-4 TBs



Compute sizing factors

Concurrency

Transaction size

Load

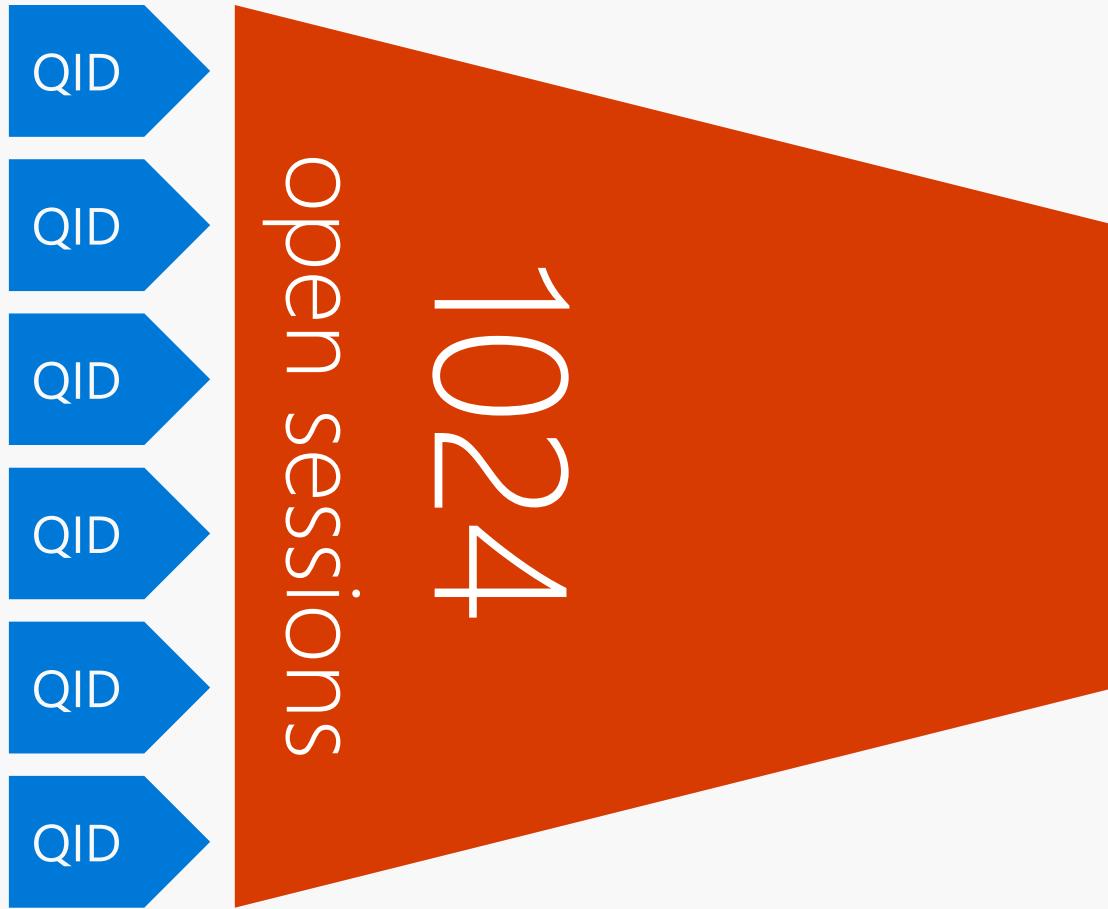
Memory



Data Warehouse Unit (DWU)

Compute sizing

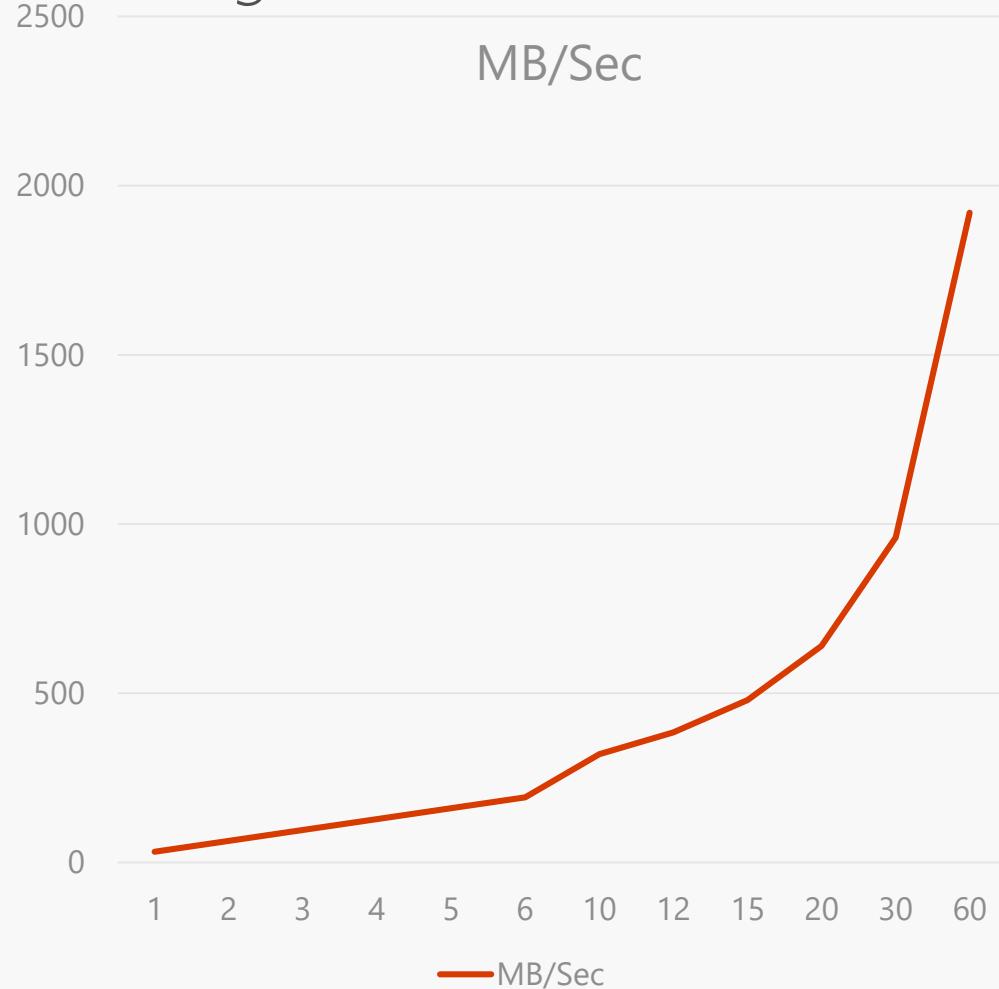
Concurrency: queries



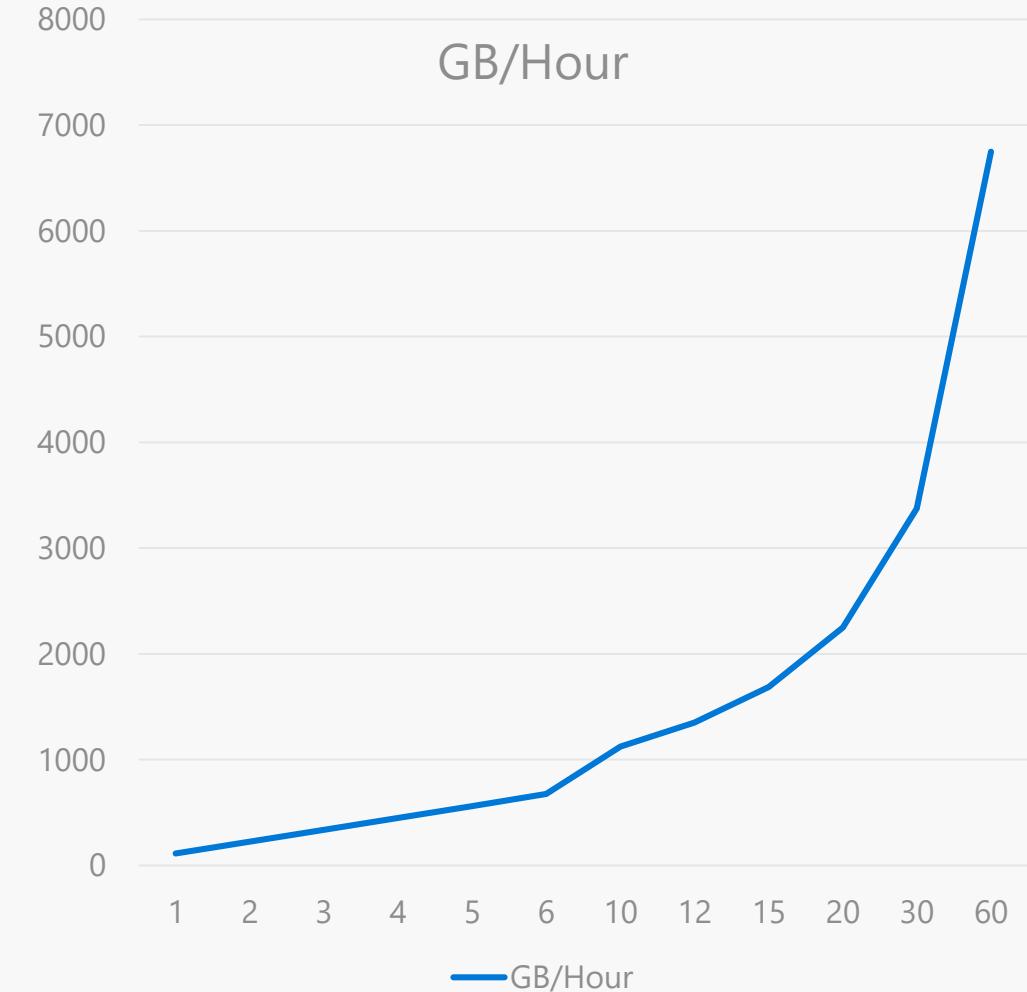
32
active queries

Compute sizing

Load scaling



30-32 MB/Sec/Node



110-115 GB/Hour/Node

Compute sizing

Loading delimited text

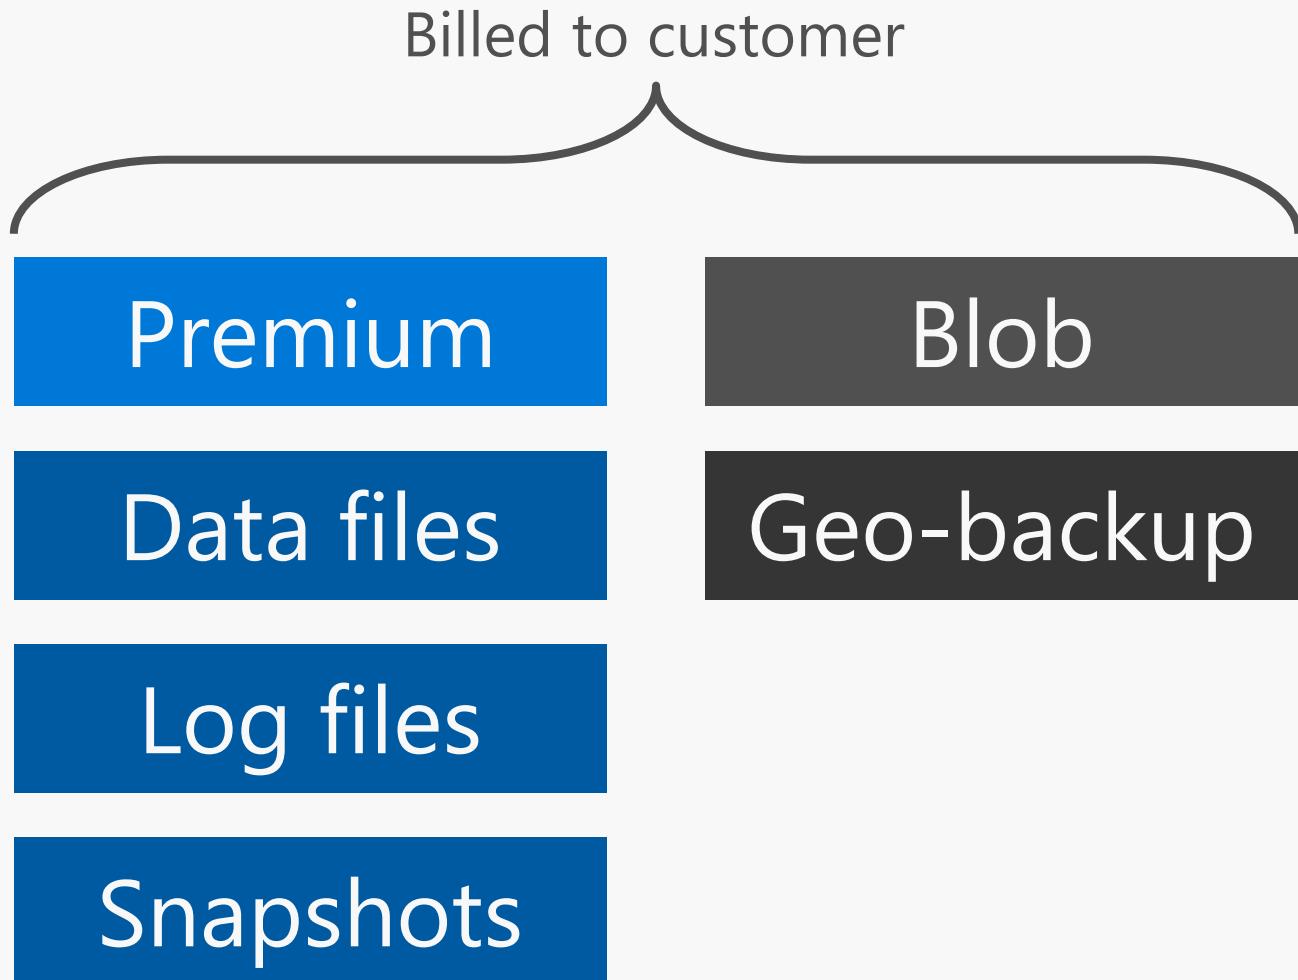
A compressed text
file cannot be read
in parallel

Splitting data across
multiple files
maximises load
performance



Storage sizing

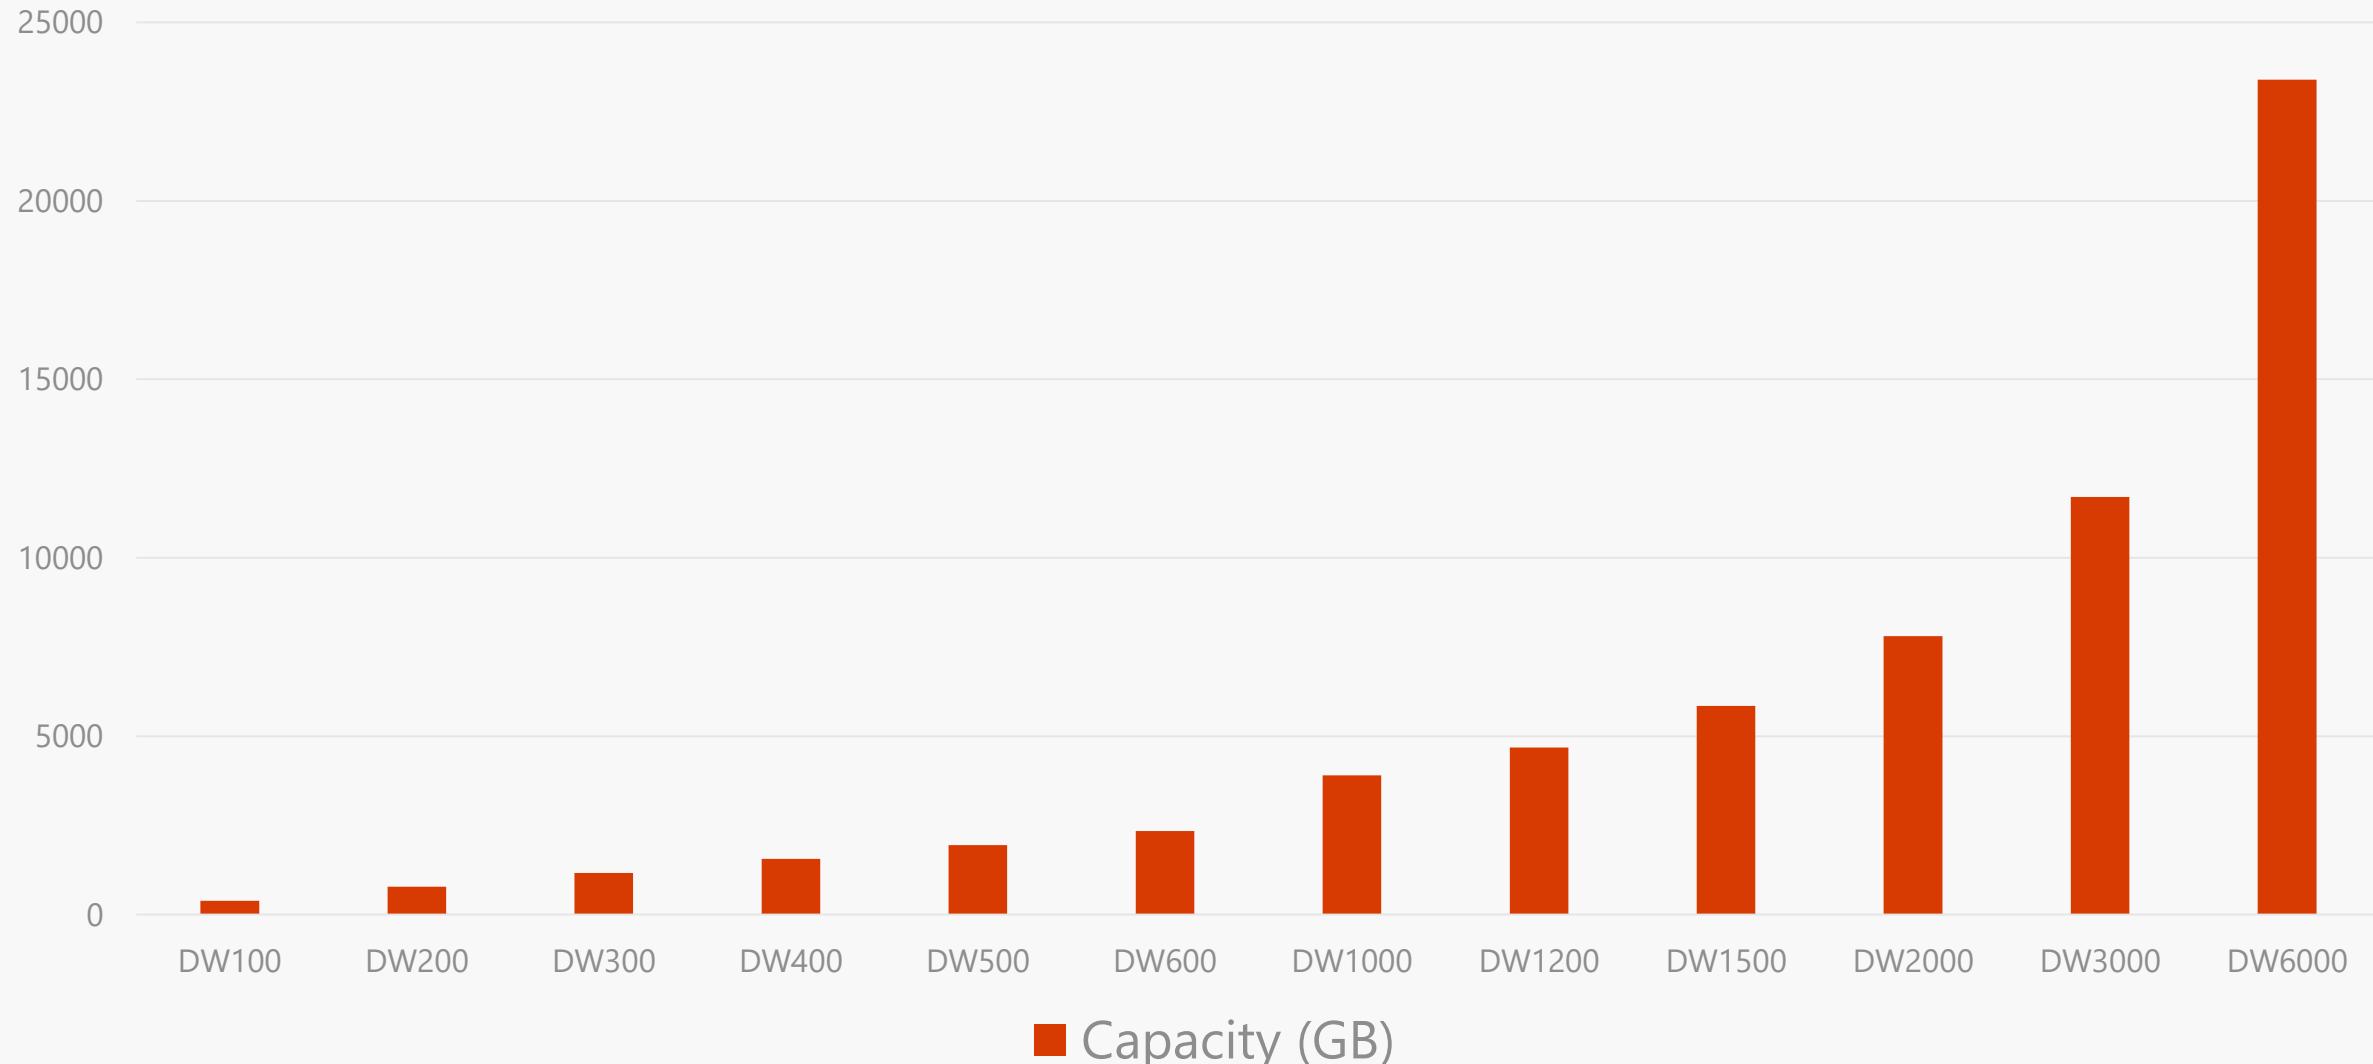
Data locations



Storage sizing

Local Storage: Temp db sizing

~399GB
per DW100



Storage sizing

Premium storage: Capacity limits

240TB

File capacity

5x

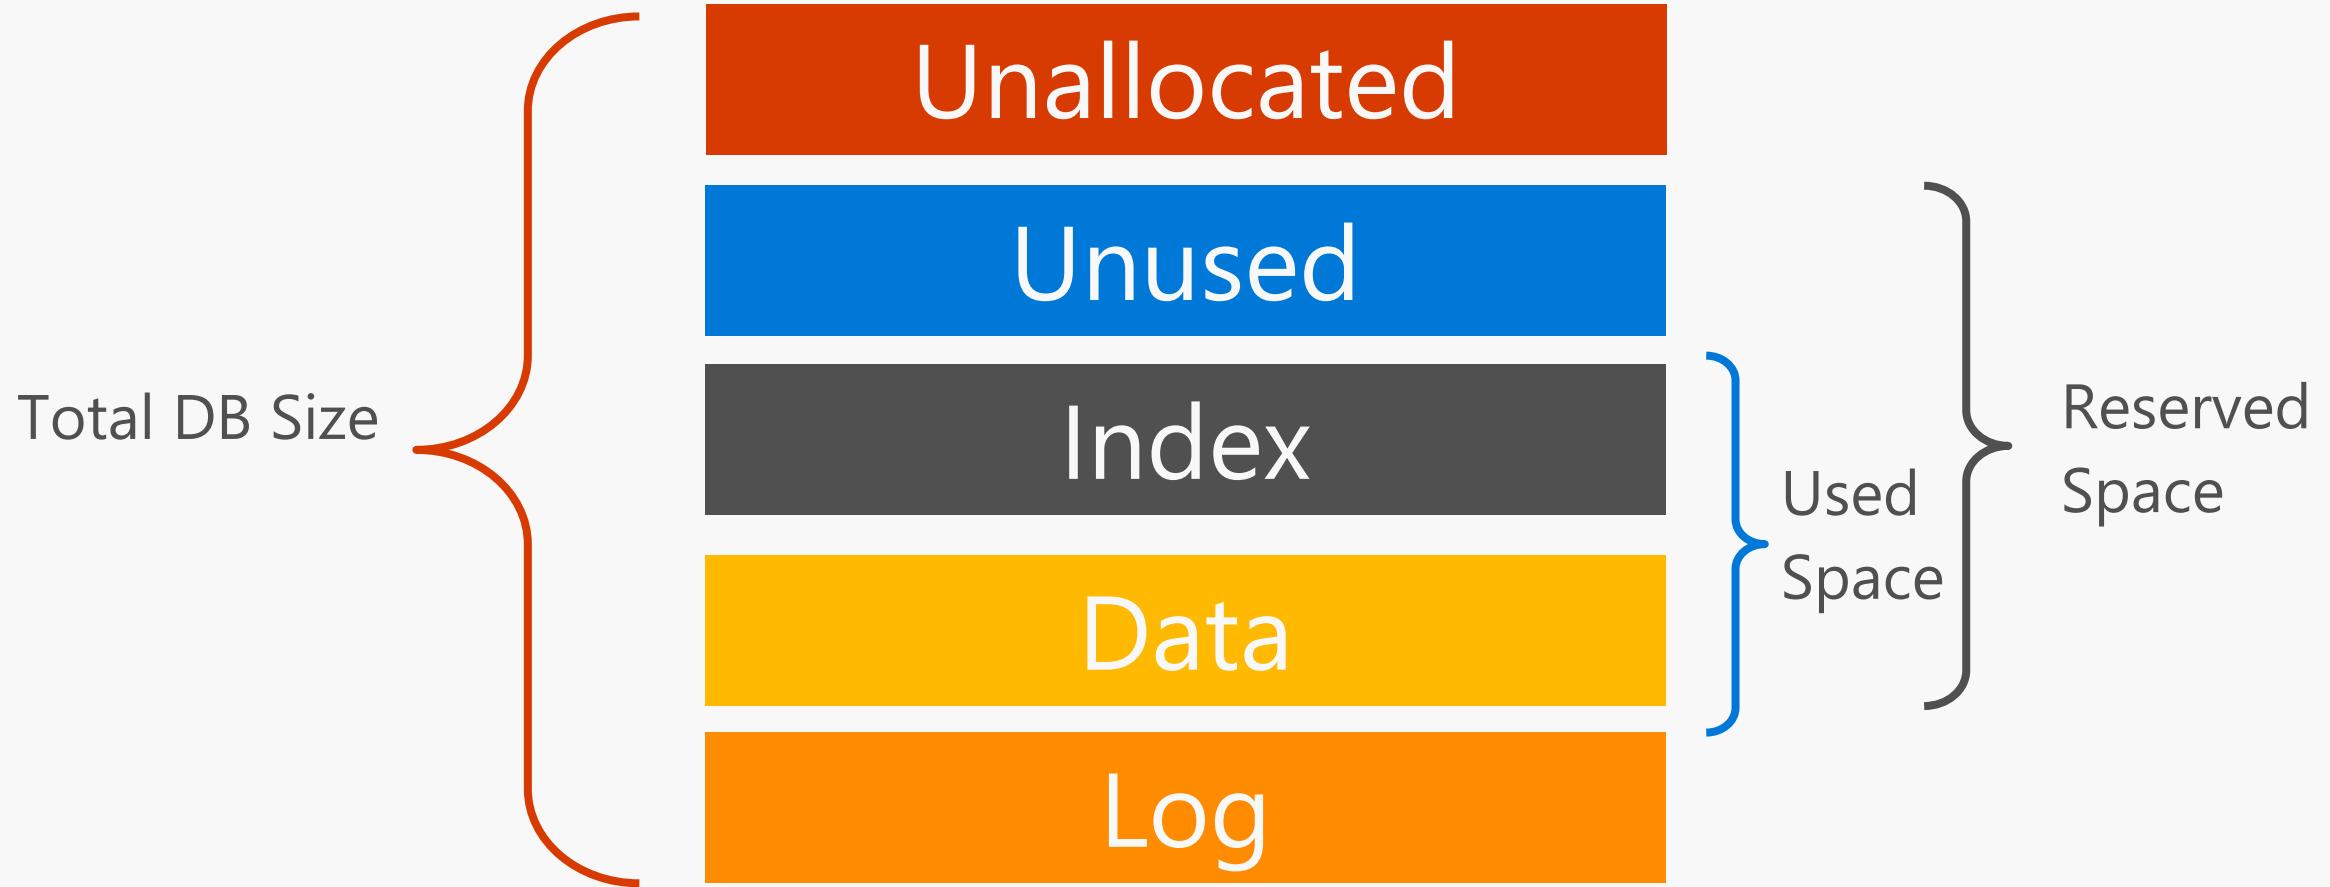
CCI compression

> 1 PB

Db capacity

Storage sizing

Premium Storage: Database Size



Storage sizing

Premium Storage: Snapshots

Frequency

Retention

4

hours

7

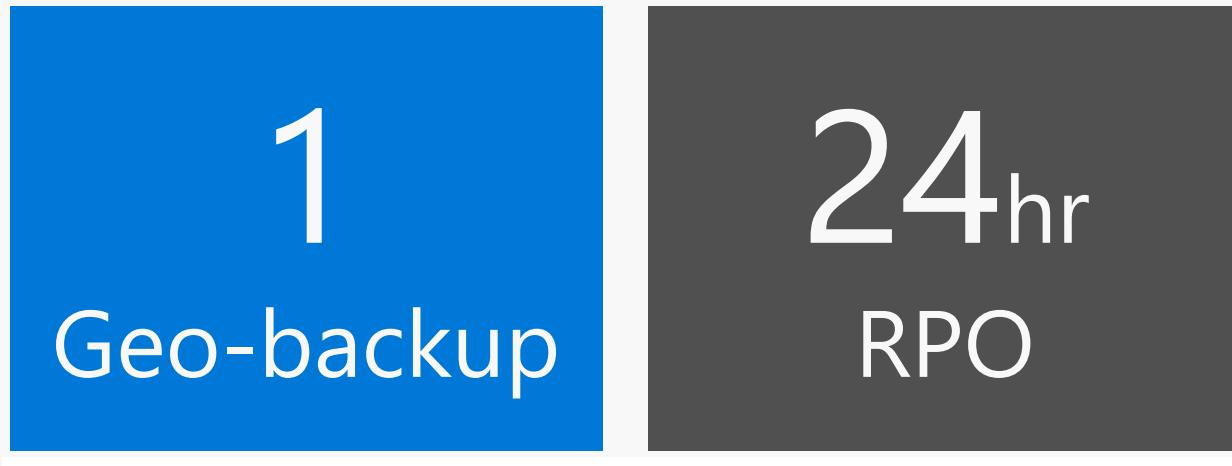
days

RPO : 8 Hours

Storage sizing

Blob Storage: Geo-redundant backups

Frequency and Retention



Geo-backup policy

Save Discard Feedback



Geo-backup copies one of the automatic snapshots each day to RA-GRS storage. This can be used in an event of a disaster to recover your data warehouse to a new region.

[Learn more](#)

Geo-backup policy

Enabled Disabled

Most recent Geo-backup time

✓ 2016-07-25T06:25:39Z

Storage sizing

Capping storage capacity

```
CREATE DATABASE MyDB COLLATE  
SQL_Latin1_General_CI_AS (  
    EDITION          = 'DataWarehouse',  
    SERVICE_OBJECTIVE = 'DW400',  
    MAXSIZE          = 10240 GB );
```

```
ALTER DATABASE MyDB  
MODIFY (MAXSIZE = 245760 GB);
```

Storage sizing

Summary

Database size:

`sp_spaceused`

Table sizing:

DMVs

Snapshot size:

Total storage size (portal) – database size (`sp_spaceused`)

Free space (unallocated):

`sp_spaceused`

Azure SQL Data Warehouse

Schema Design

Distributed Data = Buckets of Water



Evenly spreading the data leads to even use application resources

Table Distribution Options

Hash Distributed

Data divided across nodes based on hashing algorithm

Same value will always hash to same distribution

Single column only

Check for Data Skew,
NULLS, -1

Round Robin (Default)

Data distributed evenly across nodes

Easy place to start, don't need to know anything about the data

Simplicity at a cost

Will incur more data movement at query time

Replicated

(Future improvement)

Data repeated on every node

Simplifies many query plans and reduces data movement

Best with joining hash table

Consumes more space
Joining two Replicated Table runs on one node

Selecting a Distribution Method

For large fact tables, best option is to Hash Distribute

- Distribute on column that is joined to other fact tables
- Primary or surrogate key

However, be mindful of ...

- Hash column should have highly distinct values (Minimum 60 distinct values)
- Avoid distributing on a date column
- Avoid distributing on column with high frequency of NULLs and default values (e.g. -1)
- Distribution column is NOT updatable
- For compatible joins use the same data types for two distributed tables

If there are no distribution columns that make sense, then use Round Robin as last resort

Dimension Table

Small dimension table (< 60M rows)

- Clustered index
- Round Robin
- Future Improvement - Choose replicate option when it becomes available

Large dimension table

- Same design as fact table
- Clustered columnstore (by default) and distribute on join key

Optimizing with Indexes

Clustered Column Store (SQL DW Default)

- Optimal choice for large tables
- Limits scans to columns in the query
- Optimal compression
- Slower to load than Heap
- Keep partitions large enough to compress (> 1 million rows)

Heap

- Optimal choice for temporary or staging tables
- Fastest load performance

Clustered Index

- Optimal for tables < 60M rows
- Sorting operation slows down load

Non-clustered Indexes

- **Use sparingly**
- Optimize single row lookups
- Will slow down load

Partitioning

Partition on date column for archiving purposes

- Improves performance by partition elimination

Partition granularity depends on your workload

- Reload, re-process
- Aim for min 100K, optimal 1 million rows per distribution/partition (remember the 60 databases)

Optimize load performance through partition switching

Considers different grain partitions if you have hot/cold data in different tables

- Example: Hot data daily, cold data monthly

Keep the number of partitions “reasonable” as there are overheads

Re-indexing by partition when needed

DDL Example

```
CREATE TABLE FactFinance
```

```
(
```

```
    FinanceKey int NOT NULL,
```

```
    DateKey int NOT NULL,
```

```
    OrganizationKey int NOT NULL,
```

```
    DepartmentGroupKey int NOT NULL,
```

```
    ScenarioKey int NULL,
```

```
    AccountKey int NULL,
```

```
    Amount float NOT NULL)
```

```
WITH (clustered columnstore index, DISTRIBUTION = HASH(FinanceKey),
```

```
      PARTITION (DateKey RANGE RIGHT FOR VALUES
```

```
                  (20100101,20200101,20300101))
```

```
);
```

Optimizing with Statistics

Create and Update of Statistics are NOT yet automatic

- Cost Based Query Optimizer needs statistics
- Sampled stats are usually just fine
- Create statistics for all columns used in JOINs, GROUP BY, WHERE
- Update statistics after incremental load
- If needed, use multi-column statistics on join and group by

```
create statistics l_orderkey on [dbo.lineitem] (l_orderkey);
select * from sys.stats where name = 'l_orderkey';
dbcc show_statistics ("lineitem","l_orderkey");
```

Azure SQL Data Warehouse

Loading

Data Loading Options

1. PolyBase
2. BCP
3. Import/Export
4. SSIS
5. Azure Data Factory
6. Redgate Platform
7. ...

	PolyBase	BCP	SqlBulkCopy	SSIS
Rate				
Rate increase as you increase DWU	Yes	No	No	No
Rate increase as you add concurrent load	No	Yes	Yes	Yes



Best Practices – Data Loading

Polybase is the fastest method

- Upload to BLOB via AZCOPY or Powershell library
- History – use CTAS
- Incremental – use INSERT INTO

Use the highest resource class without sacrificing concurrency

Increase DWU during load, decrease when done

Known Issues:

- Do not support extended ASCII, UTF-16.
- Do not support custom multi-date format. E.g. 2000-1-6
- No reject files/reason for rejected rows.

PolyBase Characteristics

- A single PolyBase load provides best performance for non-compressed text files.
- The load performance scales as you increase DWUs.
- Parallelizes the data load process, no need to explicitly break the input data into multiple files and issue concurrent loads.
- Each reader will slice 512 MB block from data files.
- Max throughput will depend on number of readers available on the SLO you use.
- Multiple readers will not work against a compressed text file (gzip). Only a single reader is used per compressed file since uncompressing the file in the buffer is single threaded. Alternatively, generate multiple compressed files.
- The number of files should be greater than or equal to the total number of readers of your service level objective (SLO).

Loading Options

PolyBase	BCP	Import/Export	SSIS	ADF
Fastest and preferred load option. Use CTAS for initial load. Use INSERT/INTO for incremental load or CTAS into stage table and partition switch into final table.	Use only for small files < 10 GB. Limited retry logic. Does not scale as you increase DWU (single thread, single CPU on client). Increase parallel threads to improve performance.	Recommended for > 10TB. Put data on encrypted Disks and ship to Microsoft. Microsoft loads the data into SQL DW.	Increase client timeout at least 10 min, default 30 sec. Increase parallel threads to improve performance. Slight performance improvement & greater reliability if run on VM.	A simpler way to use PolyBase. Quick to configure since you don't need to define the T-SQL objects.

Azure SQL Data Warehouse

Querying

Common Data Movement Types

DMS Operation	Description
ShuffleMoveOperation	Redistributes data for compatible join or aggregation
PartitionMoveOperation	Data moves from compute to control node
BroadcastMoveOperation	Table needs to become replicated for join compatibility

Best Practices – Query Performance

Check for SKEW (DBCC PDW_SHOWSPACEUSED)

Statistics

CETAS or CTAS large return operation

De-normalize if needed

DSQL Query Plan

- Minimize data movement operations
 - Distribution & aggregation compatible
- Minimize size of data movement
 - Check for predicate pushdown. Rewrite query if needed
- Use higher resource class for memory intensive queries
- Load large external tables rather than querying directly

Azure SQL Data Warehouse

High Availability

High Availability – Out of the Box

Storage

- Backed by premium LRS
- Three copies are maintained, intra-datacenter for fault tolerance
- Replicas are served in the event of failures by the framework automatically

Compute

- Service Fabric container backbone
- Compute nodes run on containers
- No replicas maintained, new container is served up from container pool
- Less than a minute to provision
- Framework seamlessly restores

Azure SQL Data Warehouse

Backup and Disaster Recovery

Snapshots – out of the box

- Automatically scheduled to run every 4 – 8 hours
- Snapshot is incremental
- Saved to premium LRS (Locally Redundant Storage)
- Retention period is 7 days
- Point-in-time restore configurable
- Snapshots are taken only when NOT paused
- Snapshots cannot be invoked programmatically
- RPO is 8 hours
- Workarounds available to keep snapshots older than 7 days
- .Net to programmatically create snapshot
- Query to see when last snapshot started →

```
select top 1 *  
from sys.pdw_loader_backup_runs  
order by run_id desc;
```

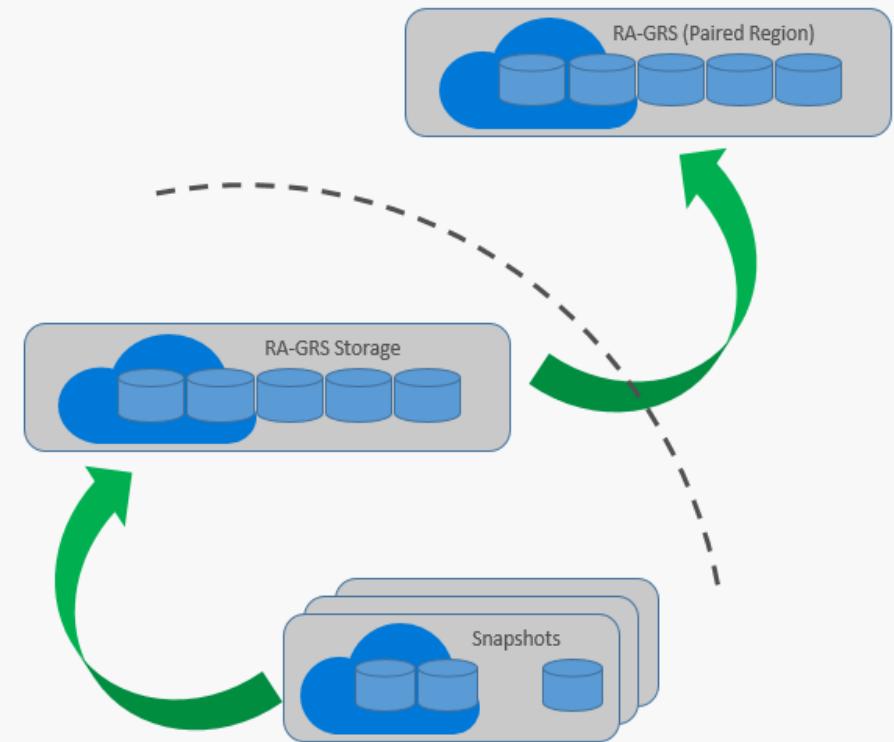
Back up to DR datacenter – out of the box

Features:

- Scheduled to run every 24 hours
- On by default, can opt-out
- RPO is 24 hours
- Full warehouse backup – only 1 backup copy
- Access your geo-backup from either the primary or the secondary location

Backup Process:

- Copies latest snapshot to a RA-GRS storage account (standard) in primary data center
- Data in this RA-GRS account gets replicated to paired Geo-DR center automatically



SQL DW – Restore (3 Methods)

Azure Portal (Method 1)

The screenshot shows the Azure Portal interface for managing a SQL Data Warehouse named XYZDWH01. The top navigation bar includes options like Settings, Pause, Scale, Open in Visual Studio, Open In PowerBI, and a redboxed Restore button. The main content area is divided into two panes: 'Essentials' and 'Monitoring'. The 'Essentials' pane displays resource group (Group), location (West Europe), and subscription name (Microsoft Azure Internal Consumption). It also shows server name (fxrdwhserver01.database.windows.net) and status (Online). The 'Monitoring' pane shows a 'Query Activity' chart with a single data series starting at 6:10 AM UTC on April 12, 2016. The right-hand pane is titled 'Restore' and contains fields for database name (XYZDWH01_2016-04-12T17-03Z), current time (2016-04-12 17:09 UTC), oldest restore point (2016-04-12 17:03 UTC), restore point (UTC) set to 2016-04-12 17:03:31 UTC, target server (fxrdwhserver01 (West Europe)), and edition (DataWarehouse (DW200)). A checkbox for 'Pin to dashboard' is checked, and a large red arrow points to the 'OK' button at the bottom right.

XYZDWH01
SQL Data Warehouse

Settings Pause Scale Open in Visual ... Open In PowerBI Restore Delete

Resource group
Group

Location
West Europe

Subscription name
Microsoft Azure Internal Consumption

Server name
fxrdwhserver01.database.windows.net

Status
Online

Connection strings
[Show database connection strings](#)

All settings →

Monitoring Add tiles **+**

Query Activity

2
1.5
1
0.5
0

Apr 12 6 AM 12 PM 6:10

There is a price implication when a new database is created.

* Database name
XYZDWH01_2016-04-12T17-03Z

Current time
2016-04-12 17:09 UTC

Oldest restore point
2016-04-12 17:03 UTC

Restore point (UTC)
2016-04-12 **17:03:31** UTC

Target server
fxrdwhserver01 (West Europe)

Edition **DataWarehouse (DW200)**

Pin to dashboard

OK

SQL DW – Restore (3 Methods)

Rest API (Method 2)

1. List all of your restorable deleted databases by using the [List restorable dropped databases](#) operation.
2. Get the details for the deleted database you want to restore by using the [Get restorable dropped database](#) operation.
3. Begin your restore by using the [Create database restore request](#) operation.
4. Track the status of your restore by using the [Database operation status](#) operation.

SQL DW – Restore (3 Methods)

Powershell (Method 3)

1. Select the subscription under your account that contains the database to be restored.
2. List restore points for the database (requires Azure Resource Management mode)
3. Pick the desired restore point using the `RestorePointCreationDate`.
4. Restore the database to the desired restore point.
5. Monitor the progress of the restore.

Azure SQL Data Warehouse

Security

Security

Authentication	Authorization	Data Protection	Connection Security
<ul style="list-style-type: none">• SQL Server authentication• Azure Active Directory• Create administrator user at provision time, then create users for SQL Server authentication	<ul style="list-style-type: none">• Same model as SQL Server – role, role memberships, permissions	<ul style="list-style-type: none">• Transparent data encryption at rest – AES 256 bit• Microsoft managed certificates and periodic rotation (90 days) by Microsoft• Encrypted configuration, automatically results in encrypted Geo-DR backup, local snapshots, metadata and log files• Encrypt using portal or T-SQL	<ul style="list-style-type: none">• Firewall rules used by both the server and the database to reject connection attempts from IP addresses that have not been explicitly whitelisted• Set up server-level firewall rule to enable client access

Auditing & Threat Detection

Auditing

- Audit log persisted in storage account
- Retain an audit trail of selected events by defining categories
- Report on database activity including some pre-configured reports.
- Excel dashboard template available
- Find suspicious events, unusual activity and trends

Threat Detection

- A new layer of security
- Detect and respond to potential threats as they occur by providing alerts on anomalous activities
- Makes it simple to address potential threats without a need for a security expert

Azure SQL Data Warehouse

Best Practices

Best practice

Cost management:

Reduce cost with pause and scale down

Pausing/scaling:

Drain transactions prior to pausing/scaling

Statistics:

Maintain statistics for performance

External tables:

Save locally, then query (versus PolyBase query)

Best practice

Scaling:

Check DMV sys.dm_pdw_exec_requests before "Change SLO" to make sure no running requests

Loading:

Use PolyBase

Scaling provides additional readers/IO

CTAS to a partition and switch partition

Use appropriate resource class for bulk insert

No singleton inserts - Batch single row inserts into larger loads

Export:

Use PolyBase

Best practice

Querying:

Keep up to date statistics

Choose distributions carefully

Size resource classes appropriately

Avoid querying external tables, load to local and then query

Use DMVs to monitor and optimize queries

Upgrading:

Use minimally logged operations for fast recovery

Build retries into applications/processes

Availability:

Build retries into applications/processes

Best practice

Partitioning:

Avoid over-partitioning – leads to poor performance

*If you issue - create 100 partitions, it translates to 100*60, 6000 partitions*

Plan for 1M rows per partition to leverage clustered columnstore

Transactions:

- Minimize transaction sizes
- Leverage special minimal logging cases, like CTAS, TRUNCATE, DROP TABLE or INSERT to empty tables, to reduce rollback risk
- For unpartitioned tables consider using a CTAS to write the data you want to keep in a table rather than using DELETE
- For partitioned tables, use Metadata Only operations like partition switching for data management

Best practice

Column-size:

Define to smallest possible size for performance

*If you issue - create 100 partitions, it translates to 100*60, 6000 partitions*

Plan for 1M rows per partition

Optimize clustered columnstore tables:

- Partitioned tables: leverage ONLY when row count > $60 * \text{number of partitions} * 1 \text{ million rows}$
- Segment quality is important (measured by number of rows in a compressed Row Group); Use User IDs in medium or large resource classes to load data; Fewer DWUs – assign higher resource class of loading user
- Columnstore tables do not yet support secondary indexes
- Queries will run faster if you select only the columns you need

Best practice

Transient data:

E.g. Stage tables

Leverage heap tables (session-access constrained) for best performance (remember to create statistics)

Use resource classes wisely:

- Resource classes are used to allocate memory to queries
- Default: all users are assigned to the small resource class which grants 100 MB of memory per distribution
- Larger resource classes, allocate more memory and therefore impact concurrency
- Certain queries, like large joins or loads to clustered columnstore tables, will benefit from larger memory allocations. Some queries, like pure scans, will see no benefit.

Best practice

Real-time ingestion and BI queries:

- Avoid real time data ingestion (ASA, Spark Streaming, Storm)
- Avoid PowerBI direct query against SQL DW
 - Use Azure Analysis Services or SSAS in a VM

Data marts:

- Use data marts to enable high concurrency workloads

Tune concurrency:

- Larger resource classes consume a higher number of concurrency slots causing other queries to queue up

Best practice

Security:

- *Restrict IP ranges when configuring server-level firewall rule*
- *Set up individual users (versus server admin)*
- *Use principle of least privilege*

Azure SQL Data Warehouse

Management & Monitoring

Management tools

Azure portal

- *Create, update, delete database*
- *Monitor database resources*

SQL Server Data Tools in Visual Studio

- *Connect, manage, develop*

Command-line tools

- *Automation; PowerShell and sqlcmd*

Dynamic Management Views

- *Bread and butter of managing SQL datawarehouse*

Monitor workloads

Dynamic Management Views:

- Monitor connections
- Monitor query execution
- Monitor queued queries
- UserID needs VIEW DATABASE STATE or CONTROL permission to query DMVs

<https://azure.microsoft.com/en-us/documentation/articles/sql-data-warehouse-manage-monitor/>

Azure SQL Data Warehouse

Migration

Why does Migration efficiency matter?

The amount of data to be migrated is large

- Usually several TB in DW to be moved from on-prem to cloud

You want to cut down the time to migrate to cloud

- Reduce downtime window
- Reduce risk of in-flight changes during migration

You want to reduce the cost of migration

- Reduce network cost
- Reduce compute cost for migration operations

Understanding the steps and tips/techniques in each helps realize an efficient migration

Data Preparation and Metadata Migration

Filter essential objects to migrate

Create performant local storage to receive exported data

Establish standard or dedicated connectivity to cloud

Chose region nearest to you with Azure SQL DW

PolyBase: One folder per table in storage container

Compatibility check must be local and as the first step

Compatibility check automation
SQL Data Warehouse Migration utility
TSQL Queries

Must complete all compatibility corrections before move

Use the SQLCAT guidance for choosing the type of distributed table

Data Migration Best Practices

Use Migration Tool - convert DDL, generate T-SQL compat report, data migration

Understand current T-SQL surface area and workarounds

Avoid Singleton DML operations (INSERT, UPDATE, DELETE)

- Batch DML if possible
- If unavoidable, wrap in transaction (BEGIN TRAN...COMMIT)

Use heap table, or temp table for “staging” data

Avoid large fully logged operations

- Considers CTAS as this is minimal logged operation
- Use LOJ as alternative for DELETE
- Process by partition to leverage parallelism and partition switching

Design retry logic to address service disruption

Data Migration Best Practices

Data Format Conversion

- Date Format, Field delimiters, escaping, field order, encoding

Compression

- Must use Gzip
- 7-Zip utility, .NET/JAVA libraries

Export

- BCP for fast export
- Multiple files per large table, one folder per table

Copy

- AZCopy
- Data Movement Library

Tips

- Incorrect format means migration needs to be entirely repeated
- Exploit bcp options, hints, parallelism
- Multiple compressed files, Split files
 - Parallel import, reliable transfer
- Don't use multiple files in the same zipfiles
- Efficient Copy
 - Parallel, Async, Resumable
 - Limit concurrent copies if low bandwidth
- Only One AzCopy per machine
- Very Large Data transfer
 - Express Route, Import/Export Service

Data Import and Transformations

Import is fastest using PolyBase

- CREATE TABLE AS...
- INSERT INTO.. SELECT FROM

Import options and Tradeoffs

- Importing Compressed
 - Faster export and transfer, slower load
- Importing Un-compressed
 - Slower export and transfer, faster load

Data Transformation

- Prefer ELT vs ETL
- Create statistics

Tips

- Increase DWUs allocated for Azure SQL DW for faster import
- Use Higher resource classes for data loading
- Refer to SQLCAT guidance on Data Loading Strategies
- Distinguish between compatibility conversions and data transformations

ETL Best Practices

Generating Identities in ETL

- Fastest done in pulling data from source
- Next best option: ETL tool generates identities

Merging into SQL DW

- Incremental load: Generate hash on columns and compare hash in DELETE, INSERT or UPDATE with SQL DW HashBytes

Sometimes ETL Tool may score over PolyBase

- Small Tables: Simpler dev experience, Error routing etc.
 - Perf Impact not noticeable
- Tables with Large Rows (>32K): No support in PolyBase yet

Tips

- Exploit SQL Server/Oracle features like ROW_NUMBER
- Use the SQL DW built in HashBytes
- Table with large sized rows and a large number of rows
- Consider vertically splitting the tables for fast loading with PolyBase

Azure SQL Data Warehouse

Pricing

<Documentation>

Azure SQL DW - Appendix

- Demo Steps
 - Create a Logical Server
 - Create a SQL DW (1000 DWU)
 - Load data (10GB) from TPCH Blob
 - Run simple queries
 - Pause and run simple queries again (Error)
 - Scale Down to 400
 - Run sample queries again
 - Walk through the code as needed for the above